

Template Demo Book: every trace on these pages was actually run

Each built-in template below is shown three ways: its real graph (drawn from the same coordinates the canvas uses), a full sim-mode run (the app's offline mode - real orchestration, canned agent text, zero model calls), and a live-mode run against the stalled-stub model from the integration tests, with the Axiom guard OFF and then ON.

Every number here was produced by an actual run of the shipped app - the same build served at nodexloop.orivael.dev - in its offline modes: seeded RNG, network disabled, no model or API key involved anywhere.

How to read the guard column

OFF -> ON compares model calls for the same stalled loop. A drop from the loop governor (LOOP_STALLED) is genuine savings. Where the drop comes from a PREFLIGHT BLOCK, the guard refused one of the template's own prompts before any call - for the prompt-injection sandbox that is the intended behaviour. Adversarial critics ('devil's advocate', 'risk judge') whose prompts legitimately say 'attack the assumptions' used to trip the same HARM pattern; they are now tagged as sanctioned personas and exempted, so they no longer false-positive. The waiver is recorded in the ledger, and a real attack - or the data flowing through the node at run time - is still classified in full. Equal counts mean the guard found nothing to stop and cost zero extra calls.

All 35 at a glance

TEMPLATE	GROUP	SHAPE	SIM RUN	LIVE CALLS OFF -> ON
blog	build	6 nodes, 7 edges, 1 loop edge(s)	7 steps, 3 iter	7 -> 7 governor saved 0
code	build	5 nodes, 5 edges, 1 loop edge(s)	7 steps, 3 iter	7 -> 7 governor saved 0
trip	build	7 nodes, 8 edges	7 steps	4 -> 4 no intervention
synth	build	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	8 -> 8 governor saved 0
landing	build	5 nodes, 6 edges, 1 loop edge(s)	7 steps, 3 iter	7 -> 7 governor saved 0
data	build	5 nodes, 5 edges, 1 loop edge(s)	5 steps	2 -> 2 no intervention
debug	build	6 nodes, 6 edges, 1 loop edge(s)	7 steps, 2 iter	3 -> 3 no intervention
rl	build	8 nodes, 9 edges, 2 loop edge(s)	11 steps, 4 iter	4 -> 4 no intervention
robot	build	7 nodes, 8 edges, 1 loop edge(s)	8 steps, 2 iter	4 -> 4 no intervention
skill	build	7 nodes, 8 edges, 2 loop edge(s)	9 steps, 4 iter	3 -> 3 no intervention
bottleneck	company	7 nodes, 8 edges, 1 loop edge(s)	9 steps, 3 iter	10 -> 8 governor saved 2
drift	company	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	11 -> 7 governor saved 4
journey	company	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	10 -> 8 governor saved 2
habit	life	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	10 -> 8 governor saved 2
decision	life	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 2 iter	9 -> 7 governor saved 2
timeaudit	life	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	10 -> 8 governor saved 2
roi	enterprise	7 nodes, 9 edges, 1 loop edge(s)	9 steps, 3 iter	10 -> 8 governor saved 2
support	enterprise	6 nodes, 6 edges, 1 loop edge(s)	8 steps, 3 iter	11 -> 7 governor saved 4
selfheal	enterprise	6 nodes, 6 edges, 1 loop edge(s)	8 steps, 3 iter	14 -> 8 governor saved 6
invoice	enterprise	6 nodes, 6 edges, 1 loop edge(s)	9 steps, 4 iter	11 -> 7 governor saved 4
certgate	governance	8 nodes, 9 edges, 1 loop edge(s)	10 steps, 3 iter	13 -> 9 governor saved 4
blastradius	governance	7 nodes, 7 edges, 1 loop edge(s)	9 steps, 3 iter	10 -> 8 governor saved 2
memfirewall	governance	6 nodes, 7 edges, 1 loop edge(s)	9 steps, 4 iter	11 -> 7 governor saved 4
supplychain	governance	7 nodes, 8 edges, 1 loop edge(s)	9 steps, 3 iter	10 -> 8 governor saved 2
costrouter	governance	7 nodes, 7 edges, 1 loop edge(s)	9 steps, 3 iter	15 -> 9 governor saved 6
factcheck	meta	5 nodes, 6 edges, 1 loop edge(s)	7 steps, 3 iter	13 -> 7 governor saved 6
injection	meta	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	13 -> 7 governor saved 6
oscillation	meta	5 nodes, 6 edges, 2 loop edge(s)	6 steps, 3 iter	24 -> 6 governor saved 18
ctxdrift	meta	5 nodes, 6 edges, 1 loop edge(s)	8 steps, 4 iter	21 -> 7 governor saved 14
handoff	meta	9 nodes, 11 edges, 1 loop edge(s)	12 steps, 4 iter	16 -> 10 governor saved 6
contract	enterprise	7 nodes, 8 edges, 1 loop edge(s)	9 steps, 3 iter	12 -> 8 governor saved 4
bisect	enterprise	7 nodes, 8 edges, 1 loop edge(s)	10 steps, 4 iter	27 -> 9 governor saved 18
roast	build	6 nodes, 7 edges, 1 loop edge(s)	8 steps, 3 iter	8 -> 8 governor saved 0
deescalate	enterprise	6 nodes, 6 edges, 1 loop edge(s)	8 steps, 3 iter	9 -> 7 governor saved 2
callroute	enterprise	13 nodes, 13 edges, 1 loop edge(s)	14 steps, 2 iter	7 -> 7 no intervention

Build

The original ten - content, code, data and research loops.

01 - BUILD - blog

Blog post factory

Research -> Write -> Critic -> Publish. Iterates until quality threshold is met.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Topic input	Received input - "User prompt + context"
Research agent	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Writer agent	Drafting content - 1,428 words - 4 sections - tone matched to brief
Critic agent	Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"
Critic agent	Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis"
Critic agent	Scored iteration 3/3 - Score: 0.90 - "Approved"
Final output	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

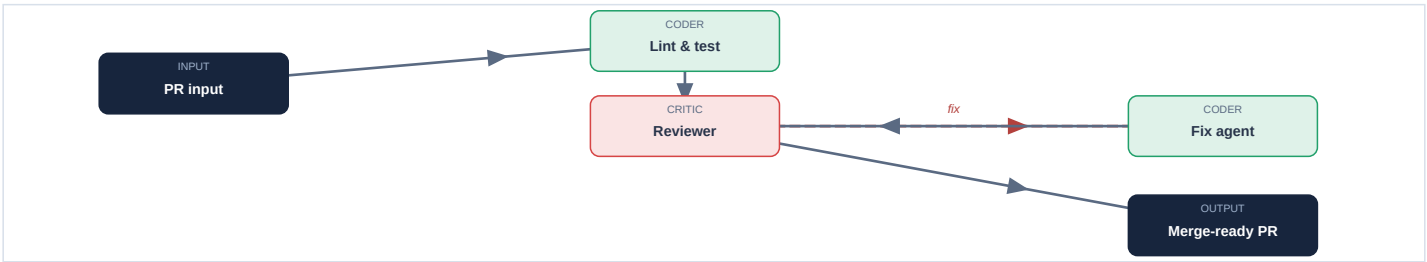
LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)
[Critic agent] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

02 - BUILD - code

Code review cycle

Lint -> Test -> Review -> Fix. Automated PR review with iterative feedback.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

PR input	Received input - "Diff + description"
Lint & test	Writing implementation - 6 files changed - lint clean - 12 tests passing
Fix agent	Writing implementation - 6 files changed - lint clean - 18 tests passing
Reviewer	Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"
Reviewer	Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis"
Reviewer	Scored iteration 3/3 - Score: 0.90 - "Approved"
Merge-ready PR	Published result - Output document ready - see the Output tab

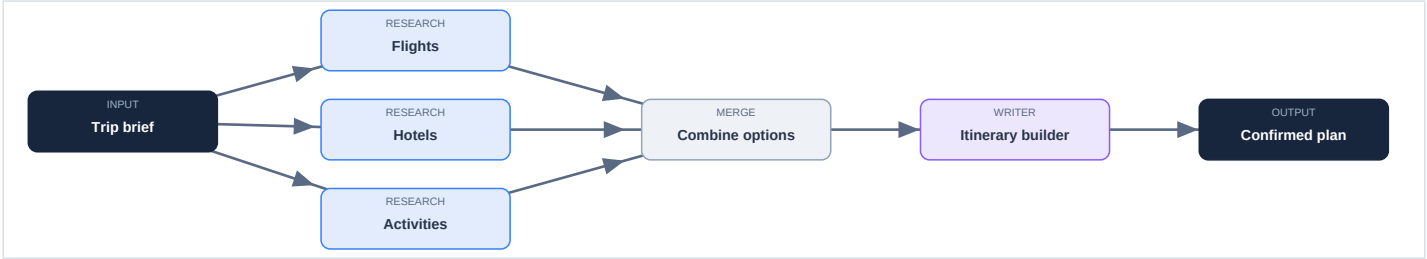
3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)
[Reviewer] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Trip planner

Search flights -> Compare hotels -> Build itinerary -> Confirm bookings.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Trip brief	Received input - "User prompt + context"
Flights	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Hotels	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Activities	Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory
Combine options	Combined streams - 3 input(s) merged into single context
Itinerary builder	Drafting content - 1,038 words - 6 sections - tone matched to brief
Confirmed plan	Published result - Output document ready - see the Output tab

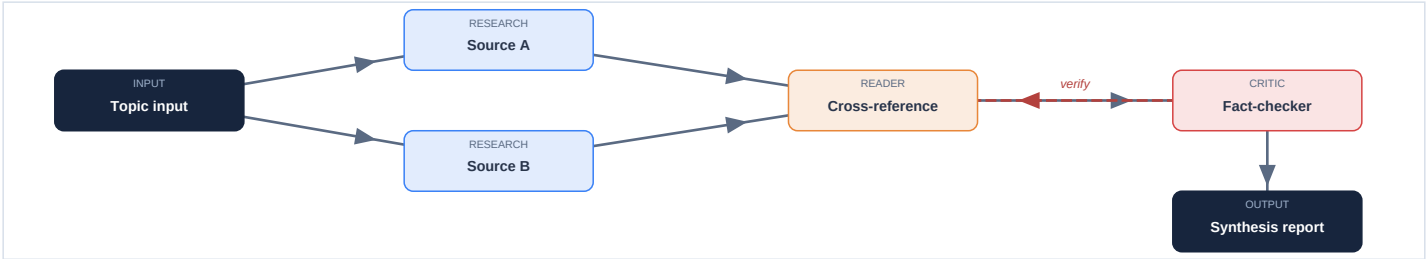
4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)
[axiom-guard] 12 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Research synthesizer

Multi-source research with cross-referencing and fact-checking loops.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Topic input	Received input - "User prompt + context"
Source A	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Source B	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Cross-reference	Reading and summarizing - 9 documents condensed - key points cross-referenced
Fact-checker	Scored iteration 1/3 - Score: 0.72 - "Claims need stronger sourcing"
Fact-checker	Scored iteration 2/3 - Score: 0.85 - "Good structure, needs tighter language"
Fact-checker	Scored iteration 3/3 - Score: 0.97 - "Approved"
Synthesis report	Published result - Output document ready - see the Output tab

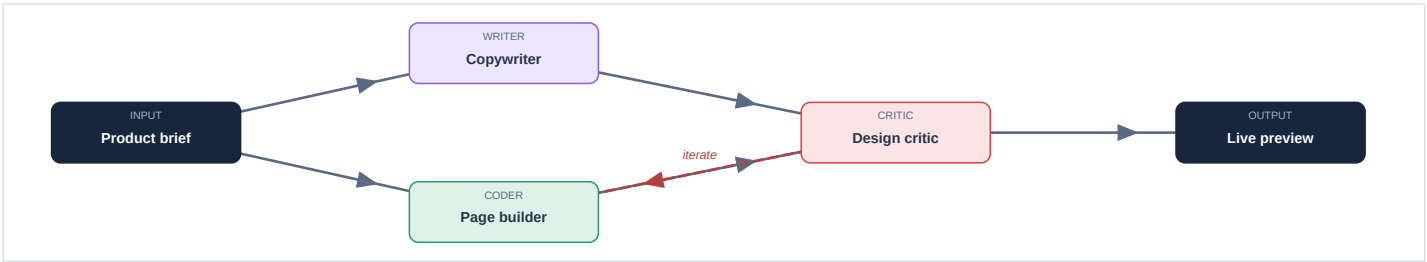
3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

8 model calls with the guard OFF -> 8 with it ON (nothing to stop; zero overhead)
[Fact-checker] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 20 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Landing page builder

Copy -> Design -> Code -> Preview. Iterates on layout and messaging.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Product brief	Received input - "User prompt + context"
Copywriter	Drafting content - 1,170 words - 6 sections - tone matched to brief
Page builder	Writing implementation - 6 files changed - lint clean - 18 tests passing
Design critic	Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"
Design critic	Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis"
Design critic	Scored iteration 3/3 - Score: 0.90 - "Approved"
Live preview	Published result - Output document ready - see the Output tab

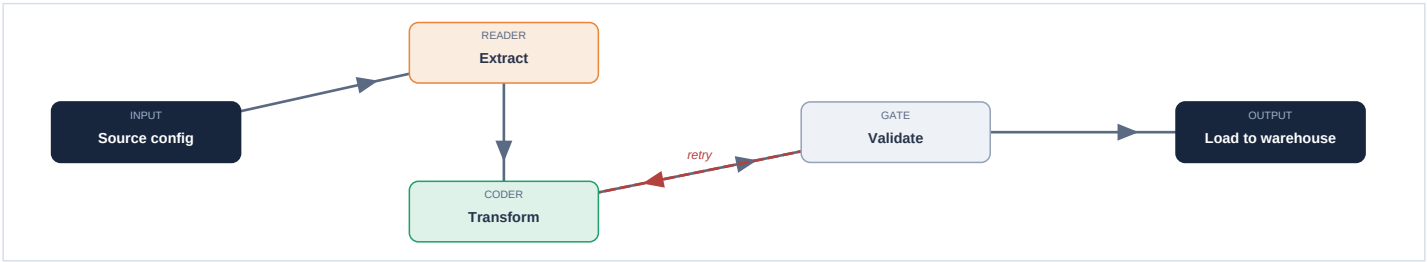
3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)
[Design critic] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Data pipeline

Extract -> Transform -> Validate -> Load. With error recovery loops.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Source config	Received input - "User prompt + context"
Extract	Reading and summarizing - 11 documents condensed - key points cross-referenced
Transform	Writing implementation - 6 files changed - lint clean - 18 tests passing
Validate	Evaluated condition - Condition true -> continue
Load to warehouse	Published result - Output document ready - see the Output tab

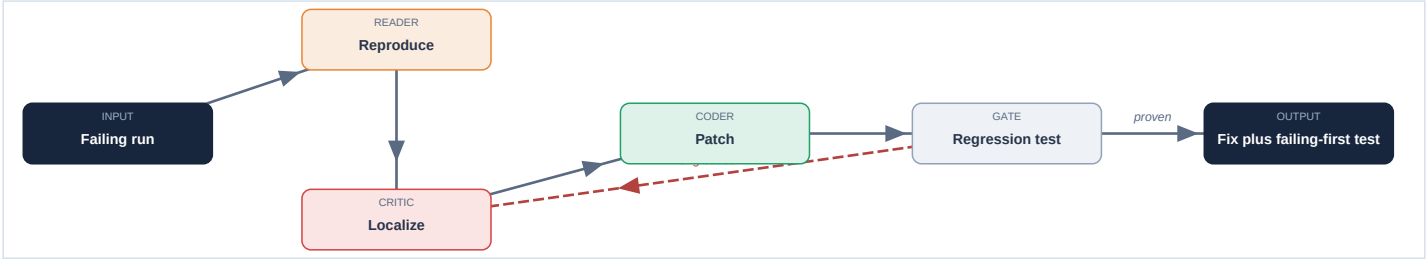
2 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

2 model calls with the guard OFF -> 2 with it ON (nothing to stop; zero overhead)
[axiom-guard] 6 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Agent failure triage

Reproduce, localize, patch, prove. Turns "the agent did something weird" into a named cause and a test that fails before the fix and passes after it.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Failing run	Received input - "Run id, transcript, the step where it went wrong"
Reproduce	Reading and summarizing - 11 documents condensed - key points cross-referenced
Localize	Scored iteration 1/3 - Score: 0.76 - "Claims need stronger sourcing"
Localize	Scored iteration 2/3 - Score: 0.88 - "Approved"
Patch	Writing implementation - 4 files changed - lint clean - 19 tests passing
Regression test	Evaluated condition - Condition true -> continue
Fix plus failing-first tes	Published result - Output document ready - see the Output tab

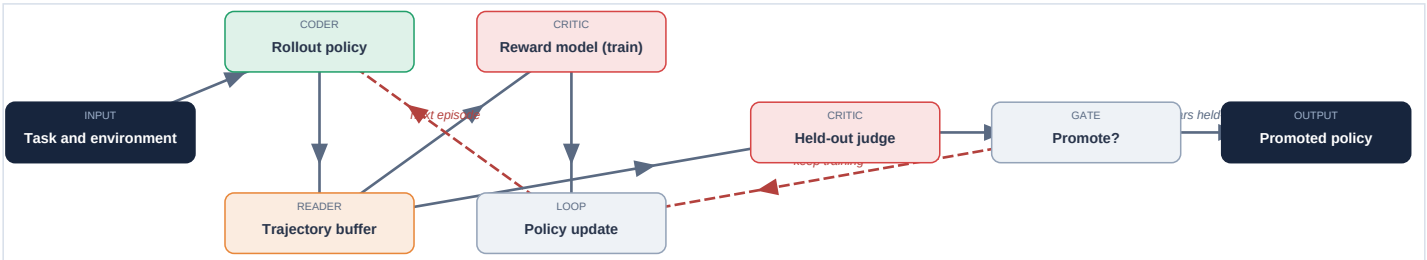
2 critic iteration(s) - final score 0.88 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

3 model calls with the guard OFF -> 3 with it ON (nothing to stop; zero overhead)
[axiom-guard] 9 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Reward loop with a held-out judge

Rollout, reward, update, evaluate. The scorer that trains the policy is never the scorer that promotes it, so reward hacking shows up as a gap between the two.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Task and environment	Received input - "Objective, action space, episode budget"
Rollout policy	Writing implementation - 6 files changed - lint clean - 12 tests passing
Trajectory buffer	Reading and summarizing - 11 documents condensed - key points cross-referenced
Reward model (train)	Scored iteration 1/8 - Score: 0.72 - "Weak transitions between key points"
... 4 more steps ...	
Held-out judge	Scored iteration 2/3 - Score: 0.86 - "Approved"
Promote?	Evaluated condition - Condition true -> continue
Promoted policy	Published result - Output document ready - see the Output tab

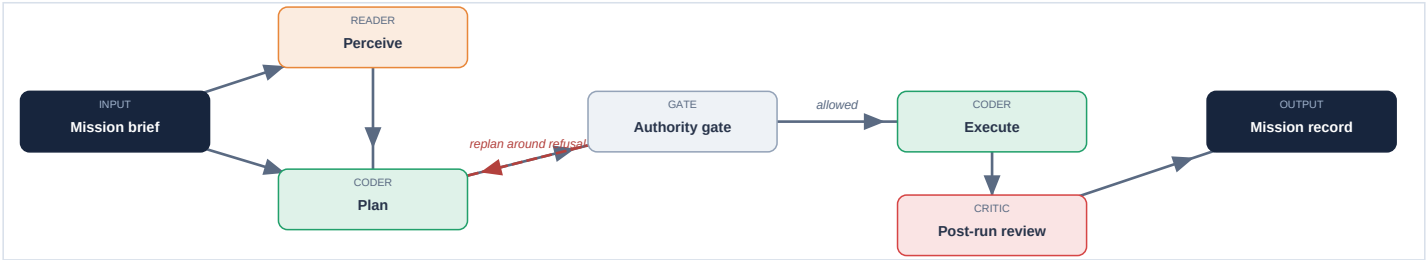
4 critic iteration(s) - final score 0.86 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)
[axiom-guard] 12 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Robot task authority

Perceive, plan, authority gate, act, review. The plan is checked against the trust level of the robot before a motor turns, and every refusal is recorded. Pairs with the Fleet sandbox.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Mission brief	Received input - "Goal, robot id, trust level, allowed zones"
Perceive	Reading and summarizing - 11 documents condensed - key points cross-referenced
Plan	Writing implementation - 6 files changed - lint clean - 18 tests passing
Authority gate	Evaluated condition - Condition true -> continue
Execute	Writing implementation - 4 files changed - lint clean - 27 tests passing
Post-run review	Scored iteration 1/3 - Score: 0.74 - "Examples are generic - add specifics"
Post-run review	Scored iteration 2/3 - Score: 0.85 - "Approved"
Mission record	Published result - Output document ready - see the Output tab

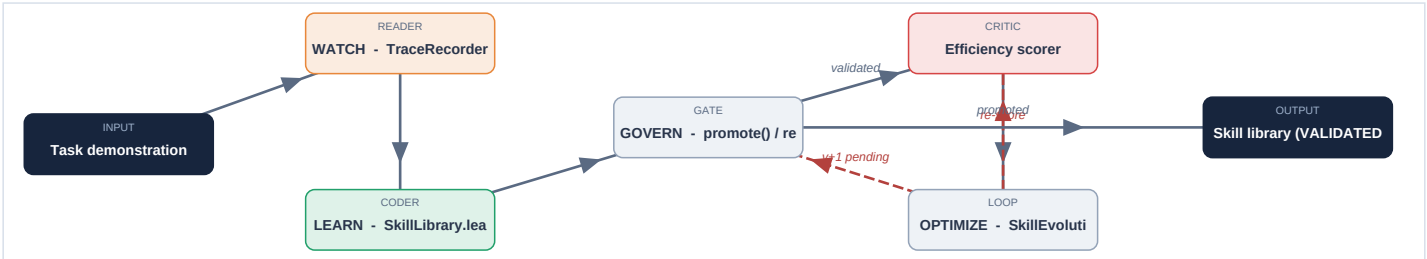
2 critic iteration(s) - final score 0.85 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)
[axiom-guard] 13 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Skill acquisition pipeline

WATCH -> LEARN -> OPTIMIZE -> GOVERN. An agent learns a reusable skill from a signed task demonstration - but nothing becomes usable until a named human promotes it. Experimental.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Task demonstration	Received input - "One real task execution - task + domain"
WATCH - TraceRecorder	Reading and summarizing - 11 documents condensed - key points cross-referenced
LEARN - SkillLibrary.lea	Writing implementation - 6 files changed - lint clean - 18 tests passing
GOVERN - promote() / rej	Evaluated condition - Condition true -> continue
... 2 more steps ...	
Efficiency scorer	Scored iteration 3/5 - Score: 0.94 - "Weak transitions between key points"
Efficiency scorer	Scored iteration 4/5 - Score: 0.99 - "Approved"
Skill library (VALIDATED o	Published result - Output document ready - see the Output tab

4 critic iteration(s) - final score 0.99 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

3 model calls with the guard OFF -> 3 with it ON (nothing to stop; zero overhead)
[axiom-guard] 9 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

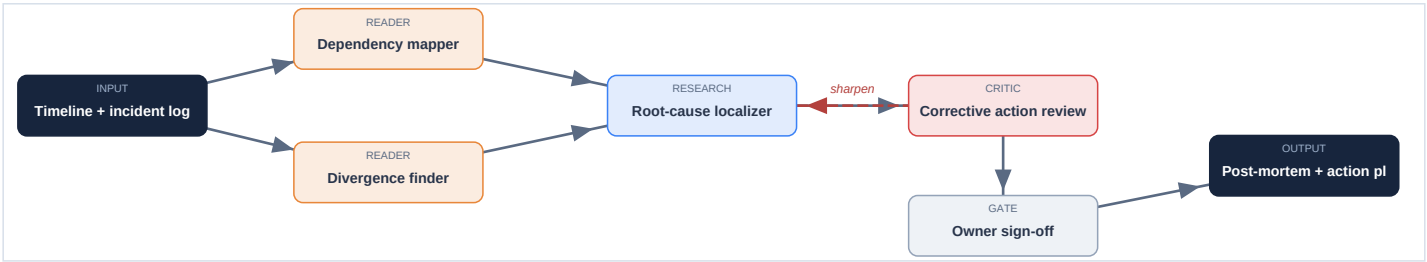
Debug your company

Org processes as loop graphs: delivery, strategy, customer journeys.

11 - COMPANY - bottleneck

Delivery bottleneck localizer

Pin the plan against what actually happened and step to the first point they diverge. Treats a stalled delivery like a hung process: find the wait, not the symptom.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Timeline + incident log	Received input - "Paste the project timeline: planned milestones, actual delivery dates, who handed off to whom, and any blockers raised."
Dependency mapper	Reading and summarizing - 11 documents condensed - key points cross-referenced
Divergence finder	Reading and summarizing - 11 documents condensed - key points cross-referenced
Root-cause localizer	Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory
... 2 more steps ...	
Corrective action reviewer	Scored iteration 3/4 - Score: 0.97 - "Approved"
Owner sign-off	Evaluated condition - Condition true -> continue
Post-mortem + action plan	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

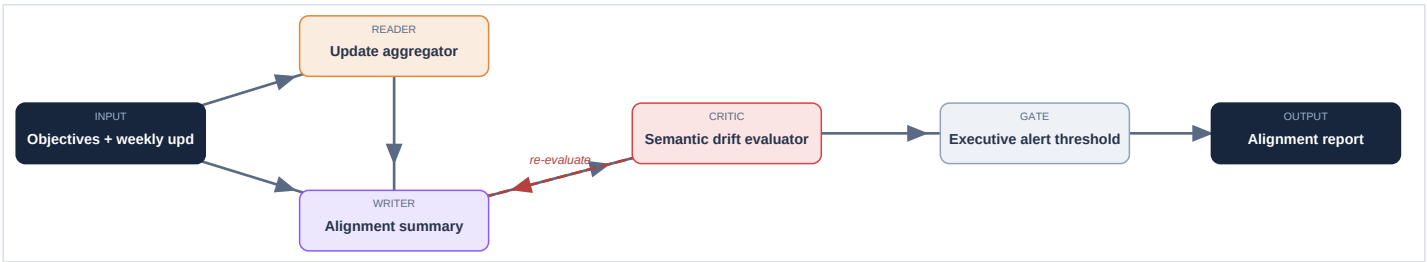
10 model calls with the guard OFF -> 8 with it ON

[Corrective action reviewer] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

12 - COMPANY - drift

Strategy drift detector

Scores what teams actually did against what the quarter said they would do. A falling score is drift - the governor sees it before the quarterly review does.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Objectives + weekly update	Received input - "Paste the quarterly objectives, then the last few weeks of team updates."
Update aggregator	Reading and summarizing - 11 documents condensed - key points cross-referenced
Alignment summary	Drafting content - 1,428 words - 4 sections - tone matched to brief
Semantic drift evaluator	Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"
Semantic drift evaluator	Scored iteration 2/5 - Score: 0.80 - "Conclusion does not follow from the analysis"
Semantic drift evaluator	Scored iteration 3/5 - Score: 0.90 - "Approved"
Executive alert threshold	Evaluated condition - Condition true -> continue
Alignment report	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

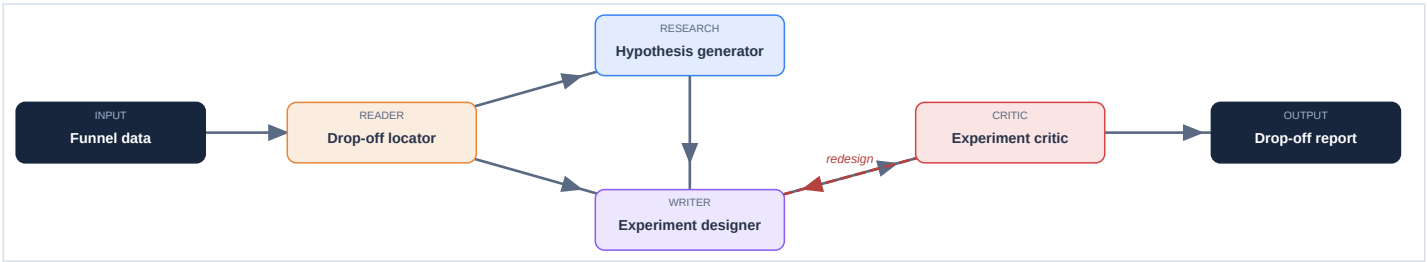
LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

11 model calls with the guard OFF -> 7 with it ON

[Semantic drift evaluator] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Customer drop-off debugger

Walk the funnel node by node to the step that loses people, generate competing explanations, and design an experiment that could actually settle it.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Funnel data	Received input - "Paste the funnel: each step, how many entered it, how many continued, and how long they spent there."
Drop-off locator	Reading and summarizing - 11 documents condensed - key points cross-referenced
Hypothesis generator	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Experiment designer	Drafting content - 1,140 words - 5 sections - tone matched to brief
Experiment critic	Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"
Experiment critic	Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"
Experiment critic	Scored iteration 3/4 - Score: 0.97 - "Approved"
Drop-off report	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

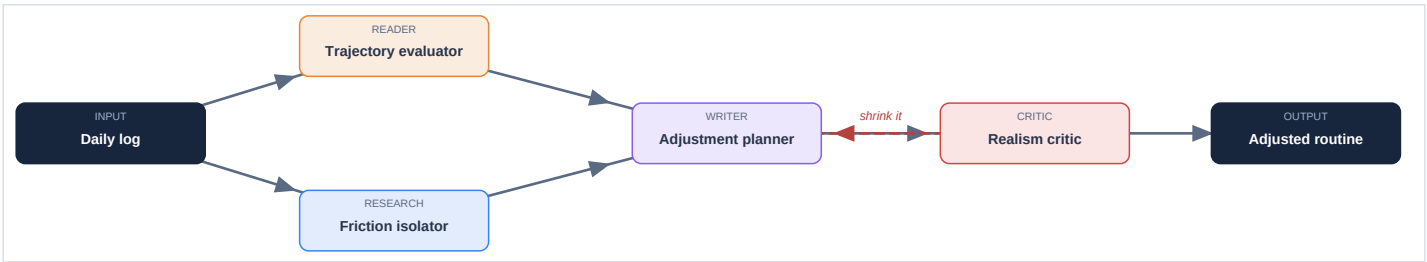
10 model calls with the guard OFF -> 8 with it ON
[Experiment critic] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 24 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Debug your life

Personal loops: habits, decisions, where the hours actually go.

Habit trajectory tracker

Why the routine collapsed in week three. Effort plotted against progress - a flat trajectory under sustained effort is the finding, and the fix lowers friction rather than demanding more discipline.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Daily log	Received input - "Paste your log: the habit you are trying to keep, which days you managed it, at what time, and what happened on the days"
Trajectory evaluator	Reading and summarizing - 11 documents condensed - key points cross-referenced
Friction isolator	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Adjustment planner	Drafting content - 1,140 words - 5 sections - tone matched to brief
Realism critic	Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"
Realism critic	Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"
Realism critic	Scored iteration 3/4 - Score: 0.97 - "Approved"
Adjusted routine	Published result - Output document ready - see the Output tab

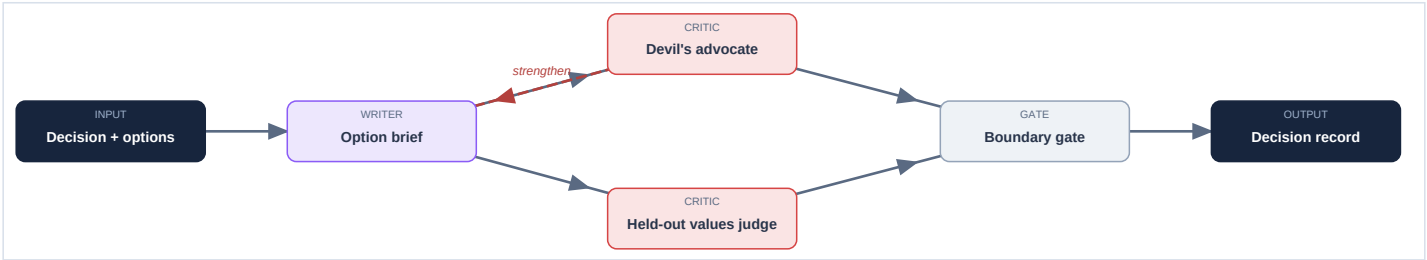
3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON
[Realism critic] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Decision preflight

Run a big decision past a devil's advocate and a held-out values judge that sits outside the loop, so it cannot be argued with. A crossed boundary stops at the gate.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Decision + options	Received input - "Describe the decision, the options you are weighing, the constraints that are real, and the values or boundaries you do "
Option brief	Drafting content - 1,170 words - 6 sections - tone matched to brief
Devil's advocate	Scored iteration 1/4 - Score: 0.76 - "Claims need stronger sourcing"
Devil's advocate	Scored iteration 2/4 - Score: 0.88 - "Approved"
Held-out values judge	Scored iteration 1/2 - Score: 0.69 - "Claims need stronger sourcing"
Held-out values judge	Scored iteration 2/2 - Score: 0.80 - "Weak transitions between key points"
Boundary gate	Evaluated condition - Condition false -> fallback path
Decision record	Published result - Output document ready - see the Output tab

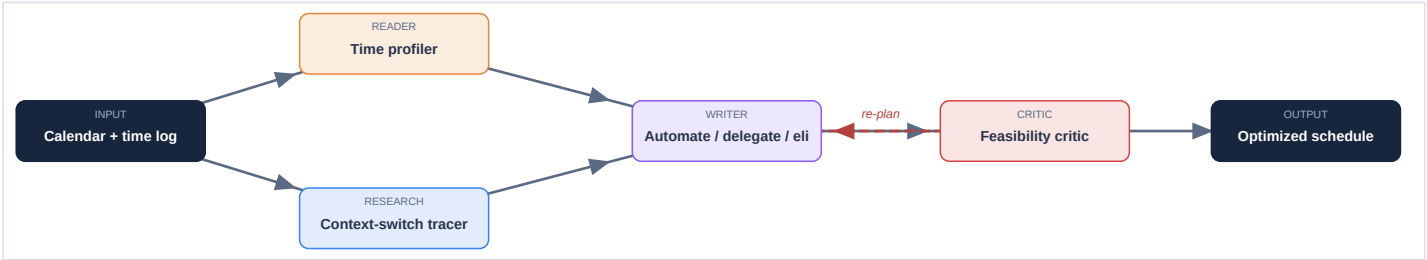
2 critic iteration(s) - final score 0.8 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

9 model calls with the guard OFF -> 7 with it ON
[Devil's advocate] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21. check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Attention leak debugger

Profile where the hours go and trace what each interruption costs downstream. Every high-drain item is automated, delegated or eliminated - keeping it is not an option.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Calendar + time log	Received input - "Paste a representative two weeks: meetings, recurring commitments, and roughly where the rest of the time went."
Time profiler	Reading and summarizing - 11 documents condensed - key points cross-referenced
Context-switch tracer	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Automate / delegate / elim	Drafting content - 1,140 words - 5 sections - tone matched to brief
Feasibility critic	Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"
Feasibility critic	Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"
Feasibility critic	Scored iteration 3/4 - Score: 0.97 - "Approved"
Optimized schedule	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON
[Feasibility critic] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21. check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

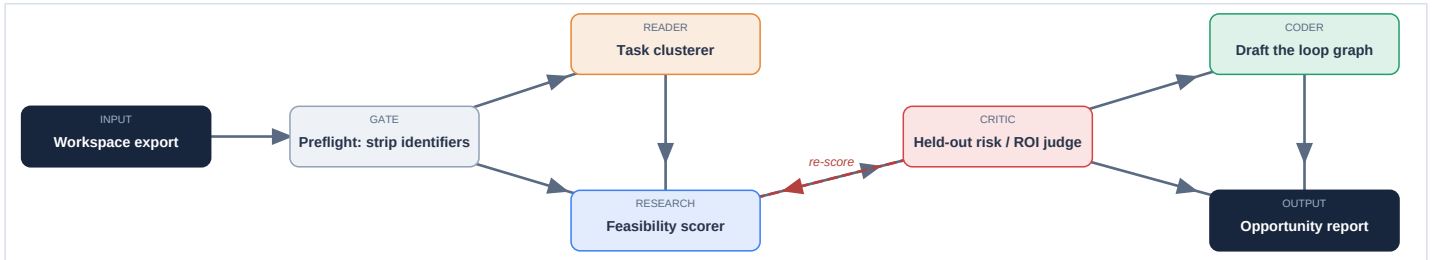
Enterprise operations

Support, self-healing CI, finance and AI-ROI auditing.

17 - ENTERPRISE - roi

AI opportunity and ROI auditor

Point it at your SOPs, tickets or wiki. It clusters the work that repeats, scores hours saved against tokens spent, has a held-out judge attack the case, and drafts a .loop.yaml for the best candidate.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Workspace export	Received input - "Paste or attach internal material: standard operating procedures, a ticket export, wiki pages, or call summaries."
Preflight: strip identifiers	Evaluated condition - Condition true -> continue
Task clusterer	Reading and summarizing - 5 documents condensed - key points cross-referenced
Feasibility scorer	Searched 7 sources - 7 sources scanned - 4 relevant - key themes extracted to shared memory
... 2 more steps ...	
Held-out risk / ROI judge	Scored iteration 3/4 - Score: 0.94 - "Approved"
Draft the loop graph	Writing implementation - 5 files changed - lint clean - 38 tests passing
Opportunity report	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.94 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON

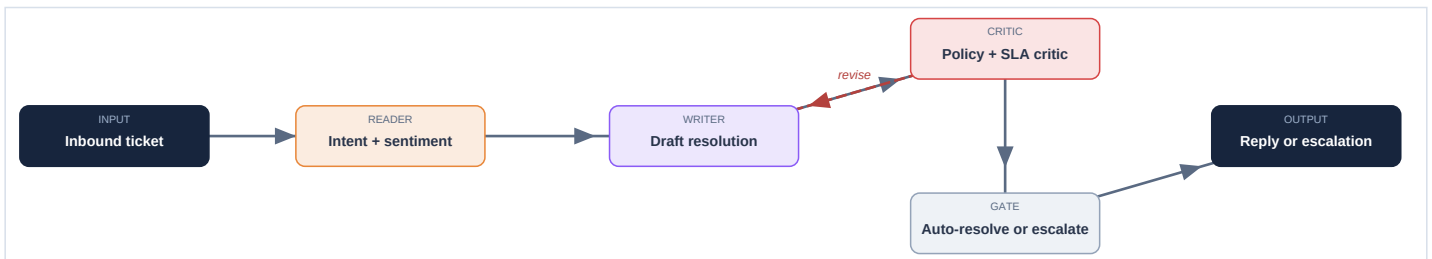
[Held-out risk / ROI judge] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 24 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

18 - ENTERPRISE - support

Support escalation loop

Classify, draft, check against policy and SLA, resolve or escalate. Catches the refund nobody authorized, and the agent rewriting the same email without adding a fix.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Inbound ticket	Received input - "Paste the ticket: what the customer said, their plan or tier, and any relevant history."
Intent + sentiment	Reading and summarizing - 11 documents condensed - key points cross-referenced
Draft resolution	Drafting content - 1,428 words - 4 sections - tone matched to brief
Policy + SLA critic	Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"
Policy + SLA critic	Scored iteration 2/5 - Score: 0.80 - "Conclusion does not follow from the analysis"
Policy + SLA critic	Scored iteration 3/5 - Score: 0.90 - "Approved"
Auto-resolve or escalate	Evaluated condition - Condition true -> continue
Reply or escalation	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

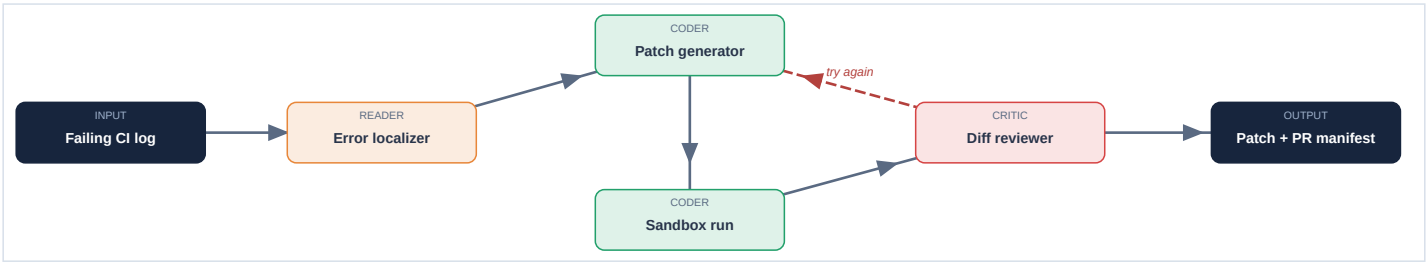
11 model calls with the guard OFF -> 7 with it ON

[Policy + SLA critic] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Self-healing CI loop

Failing build to reviewed patch: localize the first real error, fix the cause, prove it in a sandbox, and block anything that only silences the symptom.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Failing CI log	Received input - "Paste the failing CI output, the test that fails, and the change that preceded it."
Error localizer	Reading and summarizing - 11 documents condensed - key points cross-referenced
Patch generator	Writing implementation - 6 files changed - lint clean - 18 tests passing
Sandbox run	Writing implementation - 2 files changed - lint clean - 28 tests passing
Diff reviewer	Scored iteration 1/6 - Score: 0.72 - "Claims need stronger sourcing"
Diff reviewer	Scored iteration 2/6 - Score: 0.85 - "Good structure, needs tighter language"
Diff reviewer	Scored iteration 3/6 - Score: 0.97 - "Approved"
Patch + PR manifest	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

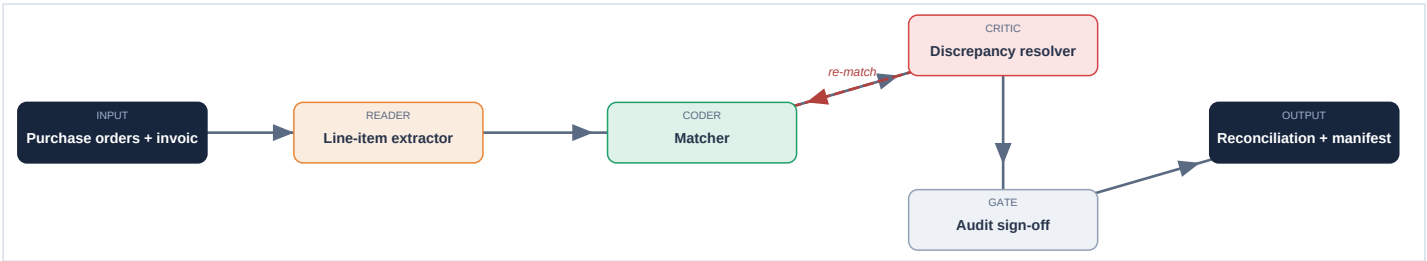
LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

14 model calls with the guard OFF -> 8 with it ON

[Diff reviewer] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Invoice reconciliation and audit

Extract line items, match both sides, explain every difference or flag it. Nothing is inferred to close a gap, and unexplained differences stop at sign-off.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Purchase orders + invoices	Received input - "Paste or attach the purchase orders and the invoices to reconcile."
Line-item extractor	Reading and summarizing - 11 documents condensed - key points cross-referenced
Matcher	Writing implementation - 6 files changed - lint clean - 18 tests passing
Discrepancy resolver	Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"
... 2 more steps ...	
Discrepancy resolver	Scored iteration 4/5 - Score: 0.98 - "Approved"
Audit sign-off	Evaluated condition - Condition true -> continue
Reconciliation + manifest	Published result - Output document ready - see the Output tab

4 critic iteration(s) - final score 0.98 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

11 model calls with the guard OFF -> 7 with it ON

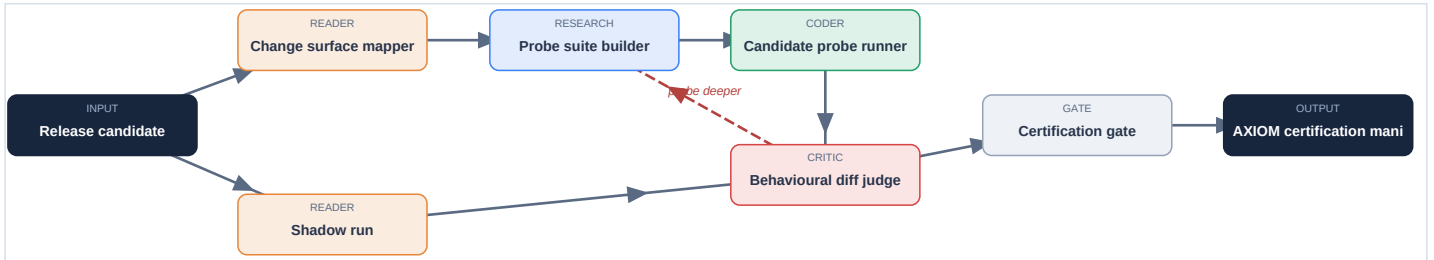
[Discrepancy resolver] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

governance

21 - GOVERNANCE - certgate

Agent release certification gate

A prompt tweak, a model upgrade or a new tool is a release. Probe the candidate against the baseline's certified behaviour, shadow-run it on real traffic, and nothing promotes except through the gate - with a version-bound certification manifest as the artifact.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Release candidate	Received input - "Describe the candidate release: what changed (prompt, model, provider, agent code, policy packs, tools), the baseline ve"
Change surface mapper	Reading and summarizing - 11 documents condensed - key points cross-referenced
Shadow run	Reading and summarizing - 11 documents condensed - key points cross-referenced
Probe suite builder	Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory
... 3 more steps ...	
Behavioural diff judge	Scored iteration 3/5 - Score: 0.96 - "Approved"
Certification gate	Evaluated condition - Condition true -> continue
AXIOM certification manife	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.96 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 9 with it ON

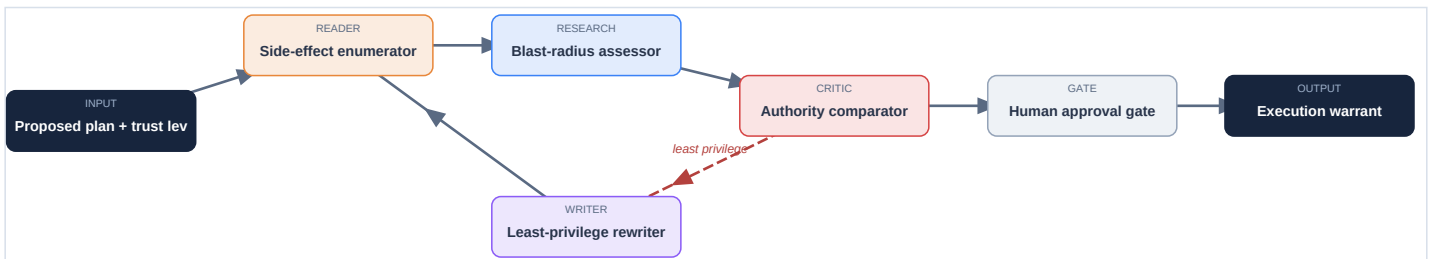
[Behavioural diff judge] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 24 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

22 - GOVERNANCE - blastradius - EXPERIMENTAL

Action blast-radius gate

Every side effect of a proposed plan - files, messages, money, deletions, credentials - enumerated and weighed against the agent's trust level before anything executes. Preview: permissions here are declared and audited, not yet enforced by the runtime.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Proposed plan + trust leve	Received input - "Paste the agent's proposed action plan before anything executes: the ordered steps, the tools each one uses, the agent's"
Least-privilege rewriter	Drafting content - 1,170 words - 6 sections - tone matched to brief
Side-effect enumerator	Reading and summarizing - 11 documents condensed - key points cross-referenced
Blast-radius assessor	Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory
... 2 more steps ...	
Authority comparator	Scored iteration 3/4 - Score: 0.97 - "Approved"
Human approval gate	Evaluated condition - Condition true -> continue
Execution warrant	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

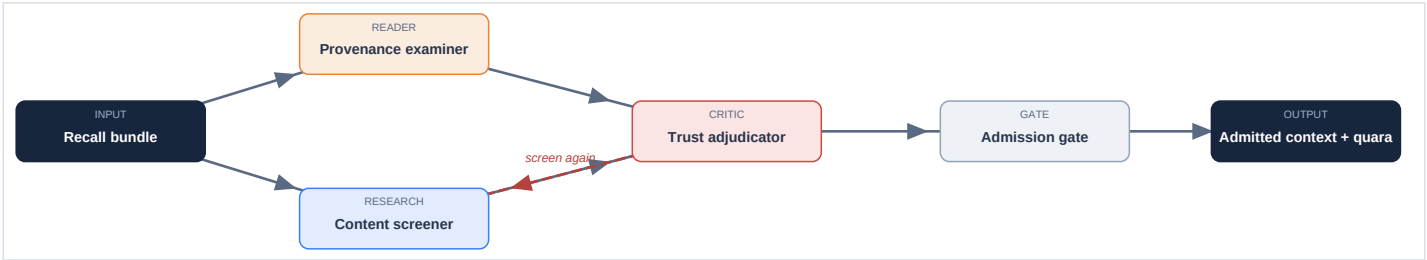
10 model calls with the guard OFF -> 8 with it ON

[Authority comparator] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Memory provenance firewall

Retrieved long-term memories are testimony, not fact. Each one is checked for origin, signature, staleness and smuggled instructions before it may enter working context - and what gets quarantined is reported, not silently dropped.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Recall bundle	Received input - "Paste the memories retrieved for this task, with whatever origin metadata and signatures they carry, plus the current ta"
Provenance examiner	Reading and summarizing - 11 documents condensed - key points cross-referenced
Content screener	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Trust adjudicator	Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"
... 2 more steps ...	
Trust adjudicator	Scored iteration 4/5 - Score: 0.98 - "Approved"
Admission gate	Evaluated condition - Condition true -> continue
Admitted context + quarant	Published result - Output document ready - see the Output tab

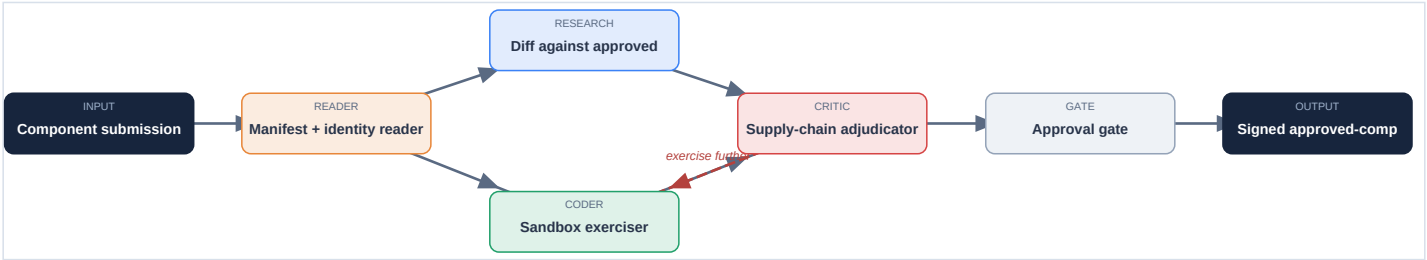
4 critic iteration(s) - final score 0.98 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

11 model calls with the guard OFF -> 7 with it ON
[Trust adjudicator] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

MCP supply-chain verifier

Before an agent, MCP server or tool joins a graph: identity, publisher, hash, permissions and schema diffed against the last approved version, then exercised in a sandbox so the verdict rests on observed behaviour. Approval is a signed manifest.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Component submission	Received input - "Paste or attach the component proposed for the graph - agent, MCP server, tool or model endpoint: its manifest, publishe"
Manifest + identity reader	Reading and summarizing - 11 documents condensed - key points cross-referenced
Diff against approved	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Sandbox exerciser	Writing implementation - 2 files changed - lint clean - 28 tests passing
... 2 more steps ...	
Supply-chain adjudicator	Scored iteration 3/4 - Score: 0.97 - "Approved"
Approval gate	Evaluated condition - Condition true -> continue
Signed approved-component	Published result - Output document ready - see the Output tab

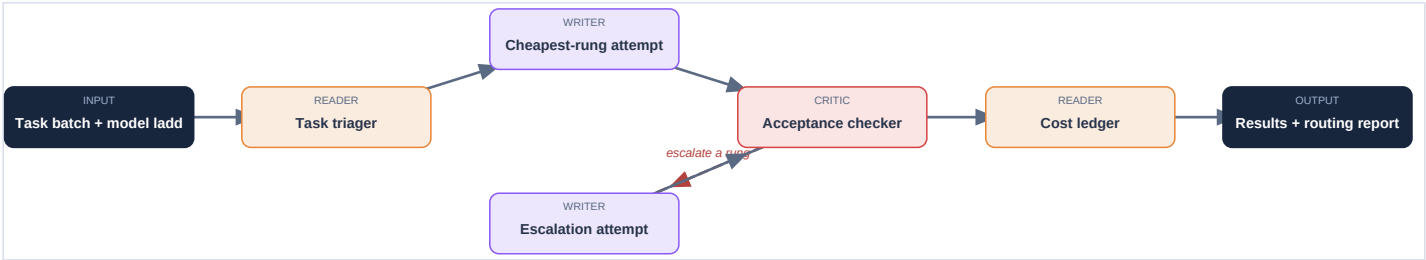
3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON
[Supply-chain adjudicator] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Cost-quality model router

The cheapest capable model goes first and a checker decides if the result actually clears the bar - escalating one rung at a time only on failure. The ledger reports cost per accepted result and savings against an all-frontier baseline, measured, not assumed.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Task batch + model ladder	Received input - "Paste the tasks to run, the model ladder from cheapest local model to frontier (add each rung as a model in Settings and"
Task triager	Reading and summarizing - 11 documents condensed - key points cross-referenced
Cheapest-rung attempt	Drafting content - 1,428 words - 4 sections - tone matched to brief
Escalation attempt	Drafting content - 1,140 words - 5 sections - tone matched to brief
... 2 more steps ...	
Acceptance checker	Scored iteration 3/6 - Score: 0.97 - "Approved"
Cost ledger	Reading and summarizing - 7 documents condensed - key points cross-referenced
Results + routing report	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

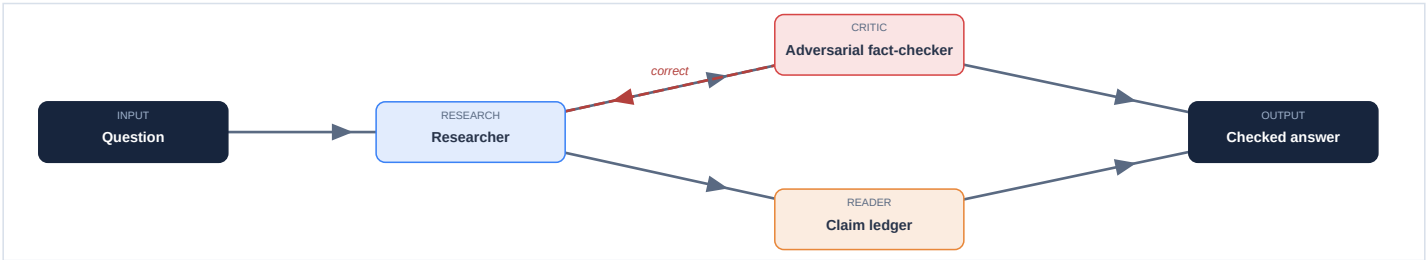
15 model calls with the guard OFF -> 9 with it ON
[Acceptance checker] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 24 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Meta

Loops that test the loops: hallucination, prompt injection, drift.

Hallucination bench

A researcher paired with a fact-checker that assumes every claim is false until the answer supports it. Break on the step where the citations start being invented.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Question	Received input - "State the question the researcher should answer, and any sources it is allowed to rely on."
Researcher	Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory
Adversarial fact-checker	Scored iteration 1/6 - Score: 0.76 - "Claims need stronger sourcing"
Adversarial fact-checker	Scored iteration 2/6 - Score: 0.88 - "Weak transitions between key points"
Adversarial fact-checker	Scored iteration 3/6 - Score: 0.97 - "Approved"
Claim ledger	Reading and summarizing - 9 documents condensed - key points cross-referenced
Checked answer	Published result - Output document ready - see the Output tab

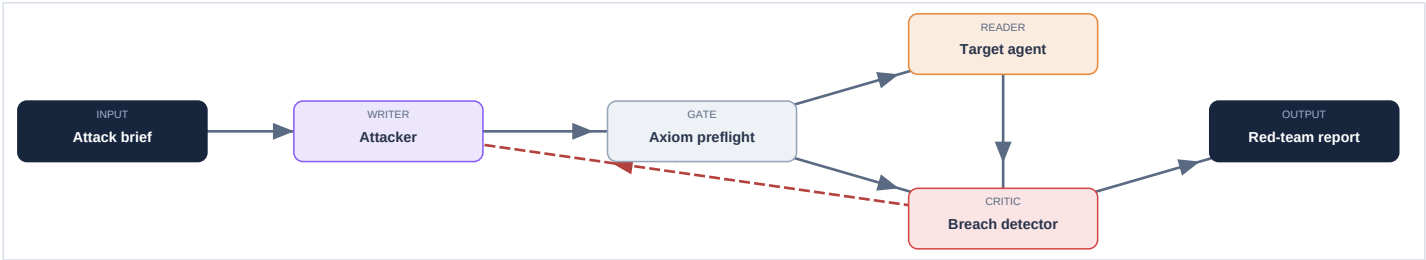
3 critic iteration(s) - final score 0.97 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 7 with it ON
[Adversarial fact-checker] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 18 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Prompt injection sandbox

An attacker loop hammering an agent with overrides, encoded payloads and indirect injection - with the local Axiom preflight in front, so a blocked attempt costs no model call at all.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Attack brief	Received input - "Describe the agent under test, the instructions it is meant to hold to, and what would count as a breach."
Attacker	Drafting content - 1,170 words - 6 sections - tone matched to brief
Axiom preflight	Evaluated condition - Condition true -> continue
Target agent	Reading and summarizing - 7 documents condensed - key points cross-referenced
Breach detector	Scored iteration 1/6 - Score: 0.71 - "Weak transitions between key points"
Breach detector	Scored iteration 2/6 - Score: 0.82 - "Conclusion does not follow from the analysis"
Breach detector	Scored iteration 3/6 - Score: 0.94 - "Approved"
Red-team report	Published result - Output document ready - see the Output tab

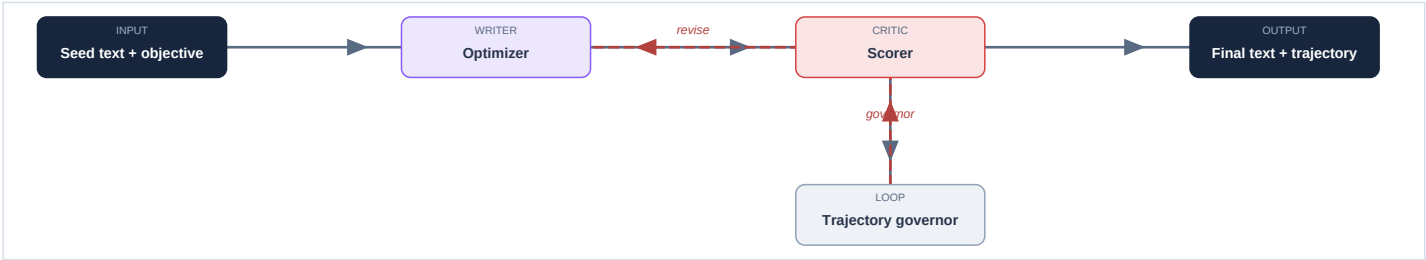
3 critic iteration(s) - final score 0.94 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 7 with it ON
[Breach detector] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 18 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Drift and oscillation detector

An optimizer and a scorer, run long enough to misbehave. Watch the trajectory flatten or turn non-monotonic, and watch the governor stop it before the budget goes.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Seed text + objective	Received input - "Paste the text to optimize and state precisely what would make it better."
Optimizer	Drafting content - 1,170 words - 6 sections - tone matched to brief
Scorer	Scored iteration 1/12 - Score: 0.76 - "Claims need stronger sourcing"
Scorer	Scored iteration 2/12 - Score: 0.88 - "Weak transitions between key points"
Scorer	Scored iteration 3/12 - Score: 0.97 - "Approved"
Final text + trajectory	Published result - Output document ready - see the Output tab

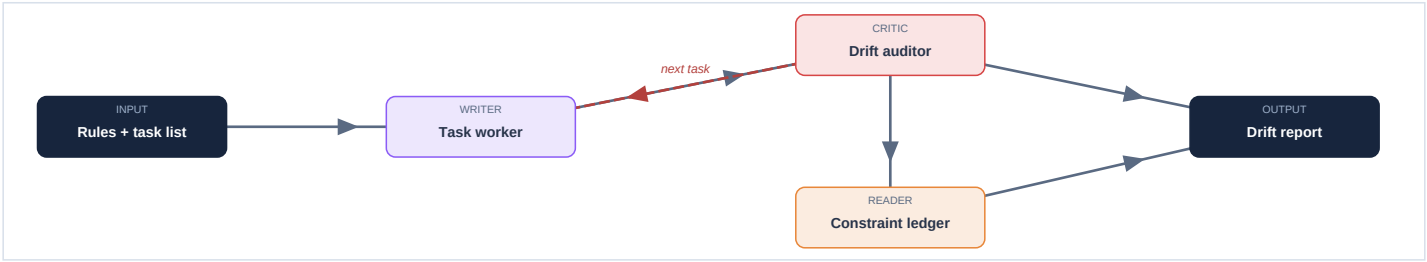
3 critic iteration(s) - final score 0.97 - 2 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

24 model calls with the guard OFF -> 6 with it ON
[Scorer] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[Trajectory governor] loop node passed context through

Context degradation stress test

Give an agent strict rules, then make it work through ten tasks. The auditor re-checks every answer against the ORIGINAL rules, so you see the iteration where the constraints start slipping rather than discovering it in production.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Rules + task list	Received input - "State the rules the agent must hold to for the whole session -- format, forbidden words, length -- then list the tasks i"
Task worker	Drafting content - 1,170 words - 6 sections - tone matched to brief
Drift auditor	Scored iteration 1/10 - Score: 0.76 - "Claims need stronger sourcing"
Drift auditor	Scored iteration 2/10 - Score: 0.88 - "Weak transitions between key points"
Drift auditor	Scored iteration 3/10 - Score: 0.97 - "Weak transitions between key points"
Drift auditor	Scored iteration 4/10 - Score: 0.99 - "Approved"
Constraint ledger	Reading and summarizing - 5 documents condensed - key points cross-referenced
Drift report	Published result - Output document ready - see the Output tab

4 critic iteration(s) - final score 0.99 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

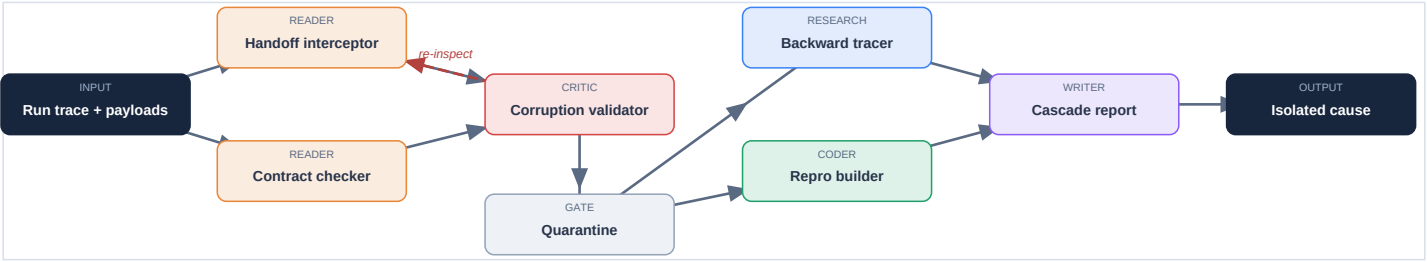
21 model calls with the guard OFF -> 7 with it ON

[Drift auditor] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 18 check(s), 0 blocked, 7 loop iteration(s) stopped early (~238 estimated tokens saved)

Handoff contagion isolator

When a downstream agent starts hallucinating, the corruption usually arrived from further up. Watches the payload at every handoff, quarantines the node that broke the cascade, and walks backwards to where the context actually went wrong.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Run trace + payloads	Received input - "Paste the run: the sequence of agents, what each one received, what each one produced, and where it went wrong."
Handoff interceptor	Reading and summarizing - 11 documents condensed - key points cross-referenced
Contract checker	Reading and summarizing - 11 documents condensed - key points cross-referenced
Corruption validator	Scored iteration 1/6 - Score: 0.72 - "Weak transitions between key points"
... 5 more steps ...	
Repro builder	Writing implementation - 2 files changed - lint clean - 32 tests passing
Cascade report	Drafting content - 1,233 words - 5 sections - tone matched to brief
Isolated cause	Published result - Output document ready - see the Output tab

4 critic iteration(s) - final score 0.98 - 6 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

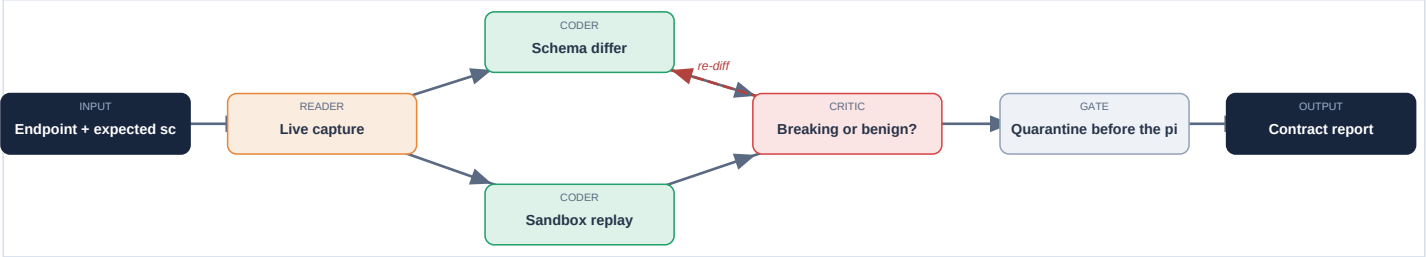
16 model calls with the guard OFF -> 10 with it ON

[Corruption validator] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 27 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Black-box contract enforcer

Treats every third-party API as a hostile input. Polls the live endpoint, diffs the real response against the schema you rely on, and traps a silent breaking change in a sandbox before it reaches your pipeline.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Endpoint + expected schema	Received input - "Give the endpoint, a sample of the response your code expects, and the fields you actually read from it."
Live capture	Reading and summarizing - 11 documents condensed - key points cross-referenced
Schema differ	Writing implementation - 6 files changed - lint clean - 18 tests passing
Sandbox replay	Writing implementation - 2 files changed - lint clean - 28 tests passing
... 2 more steps ...	
Breaking or benign?	Scored iteration 3/5 - Score: 0.97 - "Approved"
Quarantine before the pipe	Evaluated condition - Condition true -> continue
Contract report	Published result - Output document ready - see the Output tab

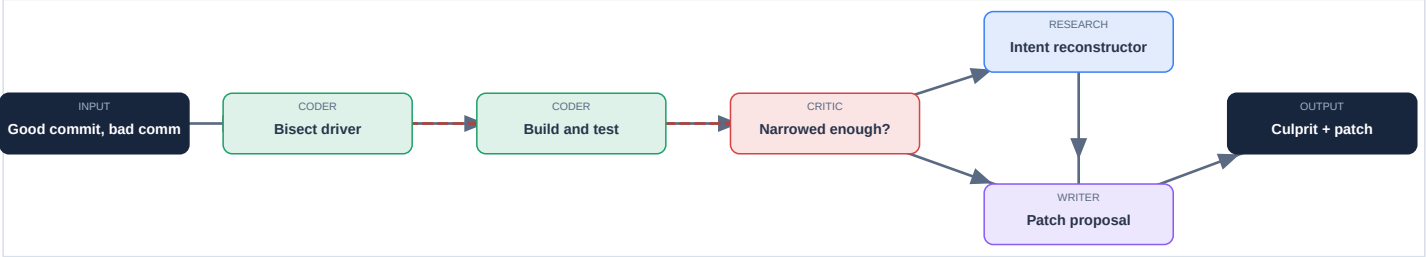
3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

12 model calls with the guard OFF -> 8 with it ON
[Breaking or benign?] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 21 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Repo regression bisector

Finds the commit that broke the build when the logs are useless, then does the part bisect cannot: reconstructs what that change was trying to do, maps it to the symptom, and drafts the patch.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Good commit, bad commit, f	Received input - "Give a commit where it worked, a commit where it does not, and the test or behaviour that fails."
Bisect driver	Writing implementation - 6 files changed - lint clean - 12 tests passing
Build and test	Writing implementation - 6 files changed - lint clean - 18 tests passing
Narrowed enough?	Scored iteration 1/12 - Score: 0.72 - "Weak transitions between key points"
... 3 more steps ...	
Intent reconstructor	Searched 12 sources - 12 sources scanned - 9 relevant - key themes extracted to shared memory
Patch proposal	Drafting content - 886 words - 5 sections - tone matched to brief
Culprit + patch	Published result - Output document ready - see the Output tab

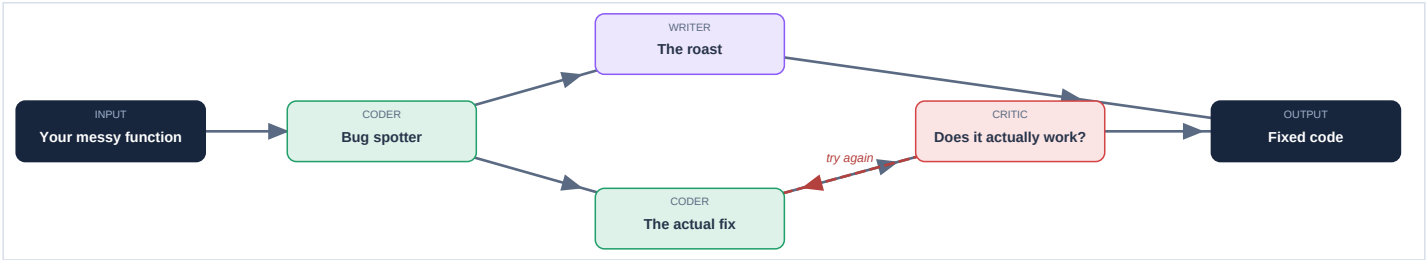
4 critic iteration(s) - final score 0.98 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

27 model calls with the guard OFF -> 9 with it ON
[Narrowed enough?] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)
[axiom-guard] 24 check(s), 0 blocked, 9 loop iteration(s) stopped early (~306 estimated tokens saved)

Sarcastic code reviewer

Drop in a messy function. It finds the real bug, hands the finding to an agent with a very different brief, and still gives you the clean fix. The shortest way to see distinct personalities handing off inside one loop.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Your messy function	Received input - "Paste a function with a bug in it. Messier is funnier."
Bug spotter	Writing implementation - 6 files changed - lint clean - 12 tests passing
The roast	Drafting content - 1,428 words - 4 sections - tone matched to brief
The actual fix	Writing implementation - 2 files changed - lint clean - 28 tests passing
Does it actually work?	Scored iteration 1/3 - Score: 0.72 - "Claims need stronger sourcing"
Does it actually work?	Scored iteration 2/3 - Score: 0.85 - "Good structure, needs tighter language"
Does it actually work?	Scored iteration 3/3 - Score: 0.97 - "Approved"
Fixed code	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

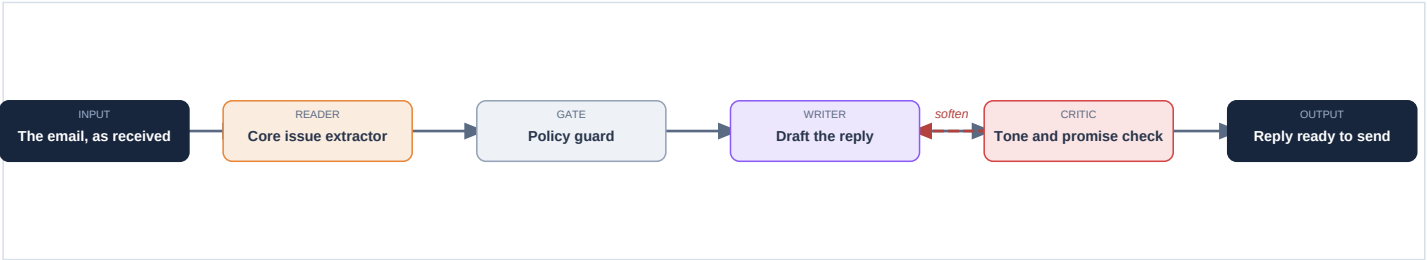
8 model calls with the guard OFF -> 8 with it ON (nothing to stop; zero overhead)

[Does it actually work?] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 21 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Angry customer de-escalator

Paste in a furious all-caps email. The loop separates the heat from the actual request, checks the proposed remedy against policy before a word is drafted, and writes a calm reply that promises only what was authorized.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

The email, as received	Received input - "Paste the customer email exactly as it arrived, plus anything relevant about their account."
Core issue extractor	Reading and summarizing - 11 documents condensed - key points cross-referenced
Policy guard	Evaluated condition - Condition true -> continue
Draft the reply	Drafting content - 1,554 words - 3 sections - tone matched to brief
Tone and promise check	Scored iteration 1/4 - Score: 0.71 - "Weak transitions between key points"
Tone and promise check	Scored iteration 2/4 - Score: 0.82 - "Conclusion does not follow from the analysis"
Tone and promise check	Scored iteration 3/4 - Score: 0.94 - "Approved"
Reply ready to send	Published result - Output document ready - see the Output tab

3 critic iteration(s) - final score 0.94 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

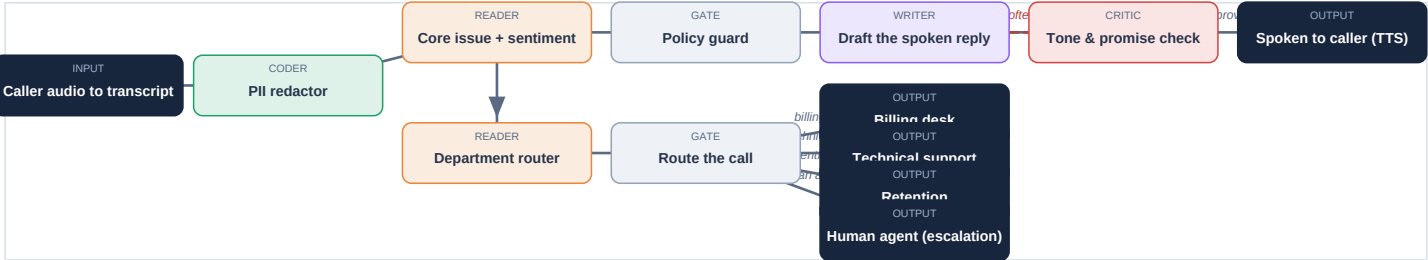
9 model calls with the guard OFF -> 7 with it ON

[Tone and promise check] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 18 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Voice de-escalation + routing

The de-escalation loop on a phone line: speech in, the calm policy-checked reply spoken back, and the call routed to the desk that owns it. A PII redactor runs first, so no model call or audit ever sees raw personal data - routing and drafting decide on topic and risk, never identity. Experimental.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

Caller audio to transcript	Received input - "Each turn arrives here as a final utterance from the speech-recognition backend (Whisper / Parakeet / Vosk), with the pr"
PII redactor	Writing implementation - 6 files changed - lint clean - 12 tests passing
Core issue + sentiment	Reading and summarizing - 11 documents condensed - key points cross-referenced
Policy guard	Evaluated condition - Condition true -> continue
... 7 more steps ...	
Technical support	Published result - Output document ready - see the Output tab
Retention	Published result - Output document ready - see the Output tab
Human agent (escalation)	Published result - Output document ready - see the Output tab

2 critic iteration(s) - final score 0.85 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)
[axiom-guard] 20 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)