

# Template Demo Book: every trace on these pages was actually run

Each built-in template below is shown three ways: its real graph (drawn from the same coordinates the canvas uses), a full sim-mode run (the app's offline mode - real orchestration, canned agent text, zero model calls), and a live-mode run against the stalled-stub model from the integration tests, with the Axiom guard OFF and then ON.

Every number here was produced by an actual run of the shipped app - the same build served at [nodexloop.orivael.dev](https://nodexloop.orivael.dev) - in its offline modes: seeded RNG, network disabled, no model or API key involved anywhere.

## How to read the guard column

OFF -> ON compares model calls for the same stalled loop. A drop from the loop governor (LOOP\_STALLED) is genuine savings. Where the drop comes from a PREFLIGHT BLOCK, the guard refused one of the template's own prompts before any call - for the prompt-injection sandbox that is the intended behaviour. Adversarial critics ('devil's advocate', 'risk judge') whose prompts legitimately say 'attack the assumptions' used to trip the same HARM pattern; they are now tagged as sanctioned personas and exempted, so they no longer false-positive. The waiver is recorded in the ledger, and a real attack - or the data flowing through the node at run time - is still classified in full. Equal counts mean the guard found nothing to stop and cost zero extra calls.

## All 34 at a glance

| TEMPLATE    | GROUP      | SHAPE                             | SIM RUN          | LIVE CALLS OFF -> ON      |
|-------------|------------|-----------------------------------|------------------|---------------------------|
| blog        | build      | 6 nodes, 7 edges, 1 loop edge(s)  | 7 steps, 3 iter  | 7 -> 7 governor saved 0   |
| code        | build      | 5 nodes, 5 edges, 1 loop edge(s)  | 7 steps, 3 iter  | 7 -> 7 governor saved 0   |
| trip        | build      | 7 nodes, 8 edges                  | 7 steps          | 4 -> 4 no intervention    |
| synth       | build      | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 8 -> 8 governor saved 0   |
| landing     | build      | 5 nodes, 6 edges, 1 loop edge(s)  | 7 steps, 3 iter  | 7 -> 7 governor saved 0   |
| data        | build      | 5 nodes, 5 edges, 1 loop edge(s)  | 5 steps          | 2 -> 2 no intervention    |
| debug       | build      | 6 nodes, 6 edges, 1 loop edge(s)  | 7 steps, 2 iter  | 3 -> 3 no intervention    |
| rl          | build      | 8 nodes, 9 edges, 2 loop edge(s)  | 11 steps, 4 iter | 4 -> 4 no intervention    |
| robot       | build      | 7 nodes, 8 edges, 1 loop edge(s)  | 8 steps, 2 iter  | 4 -> 4 no intervention    |
| skill       | build      | 7 nodes, 8 edges, 2 loop edge(s)  | 9 steps, 4 iter  | 3 -> 3 no intervention    |
| bottleneck  | company    | 7 nodes, 8 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 10 -> 8 governor saved 2  |
| drift       | company    | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 11 -> 7 governor saved 4  |
| journey     | company    | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 10 -> 8 governor saved 2  |
| habit       | life       | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 10 -> 8 governor saved 2  |
| decision    | life       | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 2 iter  | 9 -> 7 governor saved 2   |
| timeaudit   | life       | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 10 -> 8 governor saved 2  |
| roi         | enterprise | 7 nodes, 9 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 10 -> 8 governor saved 2  |
| support     | enterprise | 6 nodes, 6 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 11 -> 7 governor saved 4  |
| selfheal    | enterprise | 6 nodes, 6 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 14 -> 8 governor saved 6  |
| invoice     | enterprise | 6 nodes, 6 edges, 1 loop edge(s)  | 9 steps, 4 iter  | 11 -> 7 governor saved 4  |
| certgate    | governance | 8 nodes, 9 edges, 1 loop edge(s)  | 10 steps, 3 iter | 13 -> 9 governor saved 4  |
| blastradius | governance | 7 nodes, 7 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 10 -> 8 governor saved 2  |
| memfirewall | governance | 6 nodes, 7 edges, 1 loop edge(s)  | 9 steps, 4 iter  | 11 -> 7 governor saved 4  |
| supplychain | governance | 7 nodes, 8 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 10 -> 8 governor saved 2  |
| costrouter  | governance | 7 nodes, 7 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 15 -> 9 governor saved 6  |
| factcheck   | meta       | 5 nodes, 6 edges, 1 loop edge(s)  | 7 steps, 3 iter  | 13 -> 7 governor saved 6  |
| injection   | meta       | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 13 -> 7 governor saved 6  |
| oscillation | meta       | 5 nodes, 6 edges, 2 loop edge(s)  | 6 steps, 3 iter  | 24 -> 6 governor saved 18 |
| ctxdrift    | meta       | 5 nodes, 6 edges, 1 loop edge(s)  | 8 steps, 4 iter  | 21 -> 7 governor saved 14 |
| handoff     | meta       | 9 nodes, 11 edges, 1 loop edge(s) | 12 steps, 4 iter | 16 -> 10 governor saved 6 |
| contract    | enterprise | 7 nodes, 8 edges, 1 loop edge(s)  | 9 steps, 3 iter  | 12 -> 8 governor saved 4  |
| bisect      | enterprise | 7 nodes, 8 edges, 1 loop edge(s)  | 10 steps, 4 iter | 27 -> 9 governor saved 18 |
| roast       | build      | 6 nodes, 7 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 8 -> 8 governor saved 0   |
| deescalate  | enterprise | 6 nodes, 6 edges, 1 loop edge(s)  | 8 steps, 3 iter  | 9 -> 7 governor saved 2   |

# Build

The original ten - content, code, data and research loops.

## 01 - BUILD - blog

### Blog post factory

Research -> Write -> Critic -> Publish. Iterates until quality threshold is met.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                |                                                                                               |
|----------------|-----------------------------------------------------------------------------------------------|
| Topic input    | Received input - "User prompt + context"                                                      |
| Research agent | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory |
| Writer agent   | Drafting content - 1,428 words - 4 sections - tone matched to brief                           |
| Critic agent   | Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"                    |
| Critic agent   | Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis"           |
| Critic agent   | Scored iteration 3/3 - Score: 0.90 - "Approved"                                               |
| Final output   | Published result - Output document ready - see the Output tab                                 |

3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed


#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)  
[Critic agent] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

## 02 - BUILD - code

### Code review cycle

Lint -> Test -> Review -> Fix. Automated PR review with iterative feedback.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                |                                                                                     |
|----------------|-------------------------------------------------------------------------------------|
| PR input       | Received input - "Diff + description"                                               |
| Lint & test    | Writing implementation - 6 files changed - lint clean - 12 tests passing            |
| Fix agent      | Writing implementation - 6 files changed - lint clean - 18 tests passing            |
| Reviewer       | Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"          |
| Reviewer       | Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis" |
| Reviewer       | Scored iteration 3/3 - Score: 0.90 - "Approved"                                     |
| Merge-ready PR | Published result - Output document ready - see the Output tab                       |


3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)  
[Reviewer] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Trip planner

Search flights -> Compare hotels -> Build itinerary -> Confirm bookings.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                   |                                                                                               |
|-------------------|-----------------------------------------------------------------------------------------------|
| Trip brief        | Received input - "User prompt + context"                                                      |
| Flights           | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory |
| Hotels            | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory |
| Activities        | Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory   |
| Combine options   | Combined streams - 3 input(s) merged into single context                                      |
| Itinerary builder | Drafting content - 1,038 words - 6 sections - tone matched to brief                           |
| Confirmed plan    | Published result - Output document ready - see the Output tab                                 |


4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 12 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Research synthesizer

Multi-source research with cross-referencing and fact-checking loops.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                  |                                                                                               |
|------------------|-----------------------------------------------------------------------------------------------|
| Topic input      | Received input - "User prompt + context"                                                      |
| Source A         | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory |
| Source B         | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory |
| Cross-reference  | Reading and summarizing - 9 documents condensed - key points cross-referenced                 |
| Fact-checker     | Scored iteration 1/3 - Score: 0.72 - "Claims need stronger sourcing"                          |
| Fact-checker     | Scored iteration 2/3 - Score: 0.85 - "Good structure, needs tighter language"                 |
| Fact-checker     | Scored iteration 3/3 - Score: 0.97 - "Approved"                                               |
| Synthesis report | Published result - Output document ready - see the Output tab                                 |


3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

8 model calls with the guard OFF -> 8 with it ON (nothing to stop; zero overhead)  
[Fact-checker] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 20 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Landing page builder

Copy -> Design -> Code -> Preview. Iterates on layout and messaging.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|               |                                                                                     |
|---------------|-------------------------------------------------------------------------------------|
| Product brief | Received input - "User prompt + context"                                            |
| Copywriter    | Drafting content - 1,170 words - 6 sections - tone matched to brief                 |
| Page builder  | Writing implementation - 6 files changed - lint clean - 18 tests passing            |
| Design critic | Scored iteration 1/3 - Score: 0.72 - "Weak transitions between key points"          |
| Design critic | Scored iteration 2/3 - Score: 0.80 - "Conclusion does not follow from the analysis" |
| Design critic | Scored iteration 3/3 - Score: 0.90 - "Approved"                                     |
| Live preview  | Published result - Output document ready - see the Output tab                       |


3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

7 model calls with the guard OFF -> 7 with it ON (nothing to stop; zero overhead)  
[Design critic] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 17 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Data pipeline

Extract -> Transform -> Validate -> Load. With error recovery loops.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                   |                                                                                |
|-------------------|--------------------------------------------------------------------------------|
| Source config     | Received input - "User prompt + context"                                       |
| Extract           | Reading and summarizing - 11 documents condensed - key points cross-referenced |
| Transform         | Writing implementation - 6 files changed - lint clean - 18 tests passing       |
| Validate          | Evaluated condition - Condition true -> continue                               |
| Load to warehouse | Published result - Output document ready - see the Output tab                  |


2 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

2 model calls with the guard OFF -> 2 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 6 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Agent failure triage

Reproduce, localize, patch, prove. Turns "the agent did something weird" into a named cause and a test that fails before the fix and passes after it.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                |
|----------------------------|--------------------------------------------------------------------------------|
| Failing run                | Received input - "Run id, transcript, the step where it went wrong"            |
| Reproduce                  | Reading and summarizing - 11 documents condensed - key points cross-referenced |
| Localize                   | Scored iteration 1/3 - Score: 0.76 - "Claims need stronger sourcing"           |
| Localize                   | Scored iteration 2/3 - Score: 0.88 - "Approved"                                |
| Patch                      | Writing implementation - 4 files changed - lint clean - 19 tests passing       |
| Regression test            | Evaluated condition - Condition true -> continue                               |
| Fix plus failing-first tes | Published result - Output document ready - see the Output tab                  |


2 critic iteration(s) - final score 0.88 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

3 model calls with the guard OFF -> 3 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 9 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Reward loop with a held-out judge

Rollout, reward, update, evaluate. The scorer that trains the policy is never the scorer that promotes it, so reward hacking shows up as a gap between the two.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                      |                                                                                |
|----------------------|--------------------------------------------------------------------------------|
| Task and environment | Received input - "Objective, action space, episode budget"                     |
| Rollout policy       | Writing implementation - 6 files changed - lint clean - 12 tests passing       |
| Trajectory buffer    | Reading and summarizing - 11 documents condensed - key points cross-referenced |
| Reward model (train) | Scored iteration 1/8 - Score: 0.72 - "Weak transitions between key points"     |
| ... 4 more steps ... |                                                                                |
| Held-out judge       | Scored iteration 2/3 - Score: 0.86 - "Approved"                                |
| Promote?             | Evaluated condition - Condition true -> continue                               |
| Promoted policy      | Published result - Output document ready - see the Output tab                  |


4 critic iteration(s) - final score 0.86 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 12 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Robot task authority

Perceive, plan, authority gate, act, review. The plan is checked against the trust level of the robot before a motor turns, and every refusal is recorded. Pairs with the Fleet sandbox.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                 |                                                                                |
|-----------------|--------------------------------------------------------------------------------|
| Mission brief   | Received input - "Goal, robot id, trust level, allowed zones"                  |
| Perceive        | Reading and summarizing - 11 documents condensed - key points cross-referenced |
| Plan            | Writing implementation - 6 files changed - lint clean - 18 tests passing       |
| Authority gate  | Evaluated condition - Condition true -> continue                               |
| Execute         | Writing implementation - 4 files changed - lint clean - 27 tests passing       |
| Post-run review | Scored iteration 1/3 - Score: 0.74 - "Examples are generic - add specifics"    |
| Post-run review | Scored iteration 2/3 - Score: 0.85 - "Approved"                                |
| Mission record  | Published result - Output document ready - see the Output tab                  |


2 critic iteration(s) - final score 0.85 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

4 model calls with the guard OFF -> 4 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 13 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)

Skill acquisition pipeline

WATCH -> LEARN -> OPTIMIZE -> GOVERN. An agent learns a reusable skill from a signed task demonstration - but nothing becomes usable until a named human promotes it. Experimental.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                |
|----------------------------|--------------------------------------------------------------------------------|
| Task demonstration         | Received input - "One real task execution - task + domain"                     |
| WATCH - TraceRecorder      | Reading and summarizing - 11 documents condensed - key points cross-referenced |
| LEARN - SkillLibrary.lea   | Writing implementation - 6 files changed - lint clean - 18 tests passing       |
| GOVERN - promote() / rej   | Evaluated condition - Condition true -> continue                               |
| ... 2 more steps ...       |                                                                                |
| Efficiency scorer          | Scored iteration 3/5 - Score: 0.94 - "Weak transitions between key points"     |
| Efficiency scorer          | Scored iteration 4/5 - Score: 0.99 - "Approved"                                |
| Skill library (VALIDATED o | Published result - Output document ready - see the Output tab                  |

4 critic iteration(s) - final score 0.99 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

3 model calls with the guard OFF -> 3 with it ON (nothing to stop; zero overhead)  
[axiom-guard] 9 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)


# Debug your company

Org processes as loop graphs: delivery, strategy, customer journeys.

## 11 - COMPANY - bottleneck

### Delivery bottleneck localizer

Pin the plan against what actually happened and step to the first point they diverge. Treats a stalled delivery like a hung process: find the wait, not the symptom.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                   |                                                                                                                                            |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Timeline + incident log</b>    | Received input - "Paste the project timeline: planned milestones, actual delivery dates, who handed off to whom, and any blockers raised." |
| <b>Dependency mapper</b>          | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                             |
| <b>Divergence finder</b>          | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                             |
| <b>Root-cause localizer</b>       | Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory                                                |
| ... 2 more steps ...              |                                                                                                                                            |
| <b>Corrective action reviewer</b> | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                                            |
| <b>Owner sign-off</b>             | Evaluated condition - Condition true -> continue                                                                                           |
| <b>Post-mortem + action plan</b>  | Published result - Output document ready - see the Output tab                                                                              |

**3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed**


#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

**10 model calls with the guard OFF -> 8 with it ON**  
[Corrective action reviewer] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

## 12 - COMPANY - drift

### Strategy drift detector

Scores what teams actually did against what the quarter said they would do. A falling score is drift - the governor sees it before the quarterly review does.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                   |                                                                                             |
|-----------------------------------|---------------------------------------------------------------------------------------------|
| <b>Objectives + weekly update</b> | Received input - "Paste the quarterly objectives, then the last few weeks of team updates." |
| <b>Update aggregator</b>          | Reading and summarizing - 11 documents condensed - key points cross-referenced              |
| <b>Alignment summary</b>          | Drafting content - 1,428 words - 4 sections - tone matched to brief                         |
| <b>Semantic drift evaluator</b>   | Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"                  |
| <b>Semantic drift evaluator</b>   | Scored iteration 2/5 - Score: 0.80 - "Conclusion does not follow from the analysis"         |
| <b>Semantic drift evaluator</b>   | Scored iteration 3/5 - Score: 0.90 - "Approved"                                             |
| <b>Executive alert threshold</b>  | Evaluated condition - Condition true -> continue                                            |
| <b>Alignment report</b>           | Published result - Output document ready - see the Output tab                               |


**3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed**

#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

**11 model calls with the guard OFF -> 7 with it ON**  
[Semantic drift evaluator] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Customer drop-off debugger

Walk the funnel node by node to the step that loses people, generate competing explanations, and design an experiment that could actually settle it.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                      |                                                                                                                         |
|----------------------|-------------------------------------------------------------------------------------------------------------------------|
| Funnel data          | Received input - "Paste the funnel: each step, how many entered it, how many continued, and how long they spent there." |
| Drop-off locator     | Reading and summarizing - 11 documents condensed - key points cross-referenced                                          |
| Hypothesis generator | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                           |
| Experiment designer  | Drafting content - 1,140 words - 5 sections - tone matched to brief                                                     |
| Experiment critic    | Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"                                                    |
| Experiment critic    | Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"                                           |
| Experiment critic    | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                         |
| Drop-off report      | Published result - Output document ready - see the Output tab                                                           |

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON


10 model calls with the guard OFF -> 8 with it ON  
[Experiment critic] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 24 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Debug your life

Personal loops: habits, decisions, where the hours actually go.

Habit trajectory tracker

Why the routine collapsed in week three. Effort plotted against progress - a flat trajectory under sustained effort is the finding, and the fix lowers friction rather than demanding more discipline.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                      |                                                                                                                                             |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Daily log            | Received input - "Paste your log: the habit you are trying to keep, which days you managed it, at what time, and what happened on the days" |
| Trajectory evaluator | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Friction isolator    | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                                               |
| Adjustment planner   | Drafting content - 1,140 words - 5 sections - tone matched to brief                                                                         |
| Realism critic       | Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"                                                                        |
| Realism critic       | Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"                                                               |
| Realism critic       | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                                             |
| Adjusted routine     | Published result - Output document ready - see the Output tab                                                                               |


3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON  
[Realism critic] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Decision preflight

Run a big decision past a devil's advocate and a held-out values judge that sits outside the loop, so it cannot be argued with. A crossed boundary stops at the gate.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                       |                                                                                                                                             |
|-----------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Decision + options    | Received input - "Describe the decision, the options you are weighing, the constraints that are real, and the values or boundaries you do " |
| Option brief          | Drafting content - 1,170 words - 6 sections - tone matched to brief                                                                         |
| Devil's advocate      | Scored iteration 1/4 - Score: 0.76 - "Claims need stronger sourcing"                                                                        |
| Devil's advocate      | Scored iteration 2/4 - Score: 0.88 - "Approved"                                                                                             |
| Held-out values judge | Scored iteration 1/2 - Score: 0.69 - "Claims need stronger sourcing"                                                                        |
| Held-out values judge | Scored iteration 2/2 - Score: 0.80 - "Weak transitions between key points"                                                                  |
| Boundary gate         | Evaluated condition - Condition false -> fallback path                                                                                      |
| Decision record       | Published result - Output document ready - see the Output tab                                                                               |


2 critic iteration(s) - final score 0.8 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

9 model calls with the guard OFF -> 7 with it ON  
[Devil's advocate] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21. check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Attention leak debugger

Profile where the hours go and trace what each interruption costs downstream. Every high-drain item is automated, delegated or eliminated - keeping it is not an option.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                                                                    |
|----------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| Calendar + time log        | Received input - "Paste a representative two weeks: meetings, recurring commitments, and roughly where the rest of the time went." |
| Time profiler              | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                     |
| Context-switch tracer      | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                                      |
| Automate / delegate / elim | Drafting content - 1,140 words - 5 sections - tone matched to brief                                                                |
| Feasibility critic         | Scored iteration 1/4 - Score: 0.72 - "Claims need stronger sourcing"                                                               |
| Feasibility critic         | Scored iteration 2/4 - Score: 0.85 - "Good structure, needs tighter language"                                                      |
| Feasibility critic         | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                                    |
| Optimized schedule         | Published result - Output document ready - see the Output tab                                                                      |

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON  
[Feasibility critic] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21. check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)


## Enterprise operations

Support, self-healing CI, finance and AI-ROI auditing.

### 17 - ENTERPRISE - roi

#### AI opportunity and ROI auditor

Point it at your SOPs, tickets or wiki. It clusters the work that repeats, scores hours saved against tokens spent, has a held-out judge attack the case, and drafts a .loop.yaml for the best candidate.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                     |                                                                                                                                      |
|-------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------|
| <b>Workspace export</b>             | Received input - "Paste or attach internal material: standard operating procedures, a ticket export, wiki pages, or call summaries." |
| <b>Preflight: strip identifiers</b> | Evaluated condition - Condition true -> continue                                                                                     |
| <b>Task clusterer</b>               | Reading and summarizing - 5 documents condensed - key points cross-referenced                                                        |
| <b>Feasibility scorer</b>           | Searched 7 sources - 7 sources scanned - 4 relevant - key themes extracted to shared memory                                          |
| ... 2 more steps ...                |                                                                                                                                      |
| <b>Held-out risk / ROI judge</b>    | Scored iteration 3/4 - Score: 0.94 - "Approved"                                                                                      |
| <b>Draft the loop graph</b>         | Writing implementation - 5 files changed - lint clean - 38 tests passing                                                             |
| <b>Opportunity report</b>           | Published result - Output document ready - see the Output tab                                                                        |

**3 critic iteration(s) - final score 0.94 - 4 agent(s) - run completed**

#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON


**10 model calls with the guard OFF -> 8 with it ON**

[Held-out risk / ROI judge] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 24 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

### 18 - ENTERPRISE - support

#### Support escalation loop

Classify, draft, check against policy and SLA, resolve or escalate. Catches the refund nobody authorized, and the agent rewriting the same email without adding a fix.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                 |                                                                                                            |
|---------------------------------|------------------------------------------------------------------------------------------------------------|
| <b>Inbound ticket</b>           | Received input - "Paste the ticket: what the customer said, their plan or tier, and any relevant history." |
| <b>Intent + sentiment</b>       | Reading and summarizing - 11 documents condensed - key points cross-referenced                             |
| <b>Draft resolution</b>         | Drafting content - 1,428 words - 4 sections - tone matched to brief                                        |
| <b>Policy + SLA critic</b>      | Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"                                 |
| <b>Policy + SLA critic</b>      | Scored iteration 2/5 - Score: 0.80 - "Conclusion does not follow from the analysis"                        |
| <b>Policy + SLA critic</b>      | Scored iteration 3/5 - Score: 0.90 - "Approved"                                                            |
| <b>Auto-resolve or escalate</b> | Evaluated condition - Condition true -> continue                                                           |
| <b>Reply or escalation</b>      | Published result - Output document ready - see the Output tab                                              |

**3 critic iteration(s) - final score 0.9 - 3 agent(s) - run completed**


#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

**11 model calls with the guard OFF -> 7 with it ON**

[Policy + SLA critic] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Self-healing CI loop

Failing build to reviewed patch: localize the first real error, fix the cause, prove it in a sandbox, and block anything that only silences the symptom.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                     |                                                                                                       |
|---------------------|-------------------------------------------------------------------------------------------------------|
| Failing CI log      | Received input - "Paste the failing CI output, the test that fails, and the change that preceded it." |
| Error localizer     | Reading and summarizing - 11 documents condensed - key points cross-referenced                        |
| Patch generator     | Writing implementation - 6 files changed - lint clean - 18 tests passing                              |
| Sandbox run         | Writing implementation - 2 files changed - lint clean - 28 tests passing                              |
| Diff reviewer       | Scored iteration 1/6 - Score: 0.72 - "Claims need stronger sourcing"                                  |
| Diff reviewer       | Scored iteration 2/6 - Score: 0.85 - "Good structure, needs tighter language"                         |
| Diff reviewer       | Scored iteration 3/6 - Score: 0.97 - "Approved"                                                       |
| Patch + PR manifest | Published result - Output document ready - see the Output tab                                         |

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed


LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

14 model calls with the guard OFF -> 8 with it ON

[Diff reviewer] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Invoice reconciliation and audit

Extract line items, match both sides, explain every difference or flag it. Nothing is inferred to close a gap, and unexplained differences stop at sign-off.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                       |
|----------------------------|---------------------------------------------------------------------------------------|
| Purchase orders + invoices | Received input - "Paste or attach the purchase orders and the invoices to reconcile." |
| Line-item extractor        | Reading and summarizing - 11 documents condensed - key points cross-referenced        |
| Matcher                    | Writing implementation - 6 files changed - lint clean - 18 tests passing              |
| Discrepancy resolver       | Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"            |
| ... 2 more steps ...       |                                                                                       |
| Discrepancy resolver       | Scored iteration 4/5 - Score: 0.98 - "Approved"                                       |
| Audit sign-off             | Evaluated condition - Condition true -> continue                                      |
| Reconciliation + manifest  | Published result - Output document ready - see the Output tab                         |

4 critic iteration(s) - final score 0.98 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

11 model calls with the guard OFF -> 7 with it ON


[Discrepancy resolver] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

## governance

### 21 - GOVERNANCE - certgate

#### Agent release certification gate

A prompt tweak, a model upgrade or a new tool is a release. Probe the candidate against the baseline's certified behaviour, shadow-run it on real traffic, and nothing promotes except through the gate - with a version-bound certification manifest as the artifact.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                                                                             |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Release candidate          | Received input - "Describe the candidate release: what changed (prompt, model, provider, agent code, policy packs, tools), the baseline ve" |
| Change surface mapper      | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Shadow run                 | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Probe suite builder        | Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory                                                 |
| ... 3 more steps ...       |                                                                                                                                             |
| Behavioural diff judge     | Scored iteration 3/5 - Score: 0.96 - "Approved"                                                                                             |
| Certification gate         | Evaluated condition - Condition true -> continue                                                                                            |
| AXIOM certification manife | Published result - Output document ready - see the Output tab                                                                               |

3 critic iteration(s) - final score 0.96 - 5 agent(s) - run completed

#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 9 with it ON


[Behavioural diff judge] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 24 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

### 22 - GOVERNANCE - blastradius - EXPERIMENTAL

#### Action blast-radius gate

Every side effect of a proposed plan - files, messages, money, deletions, credentials - enumerated and weighed against the agent's trust level before anything executes. Preview: permissions here are declared and audited, not yet enforced by the runtime.



#### SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                                                                             |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Proposed plan + trust leve | Received input - "Paste the agent's proposed action plan before anything executes: the ordered steps, the tools each one uses, the agent's" |
| Least-privilege rewriter   | Drafting content - 1,170 words - 6 sections - tone matched to brief                                                                         |
| Side-effect enumerator     | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Blast-radius assessor      | Searched 9 sources - 9 sources scanned - 6 relevant - key themes extracted to shared memory                                                 |
| ... 2 more steps ...       |                                                                                                                                             |
| Authority comparator       | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                                             |
| Human approval gate        | Evaluated condition - Condition true -> continue                                                                                            |
| Execution warrant          | Published result - Output document ready - see the Output tab                                                                               |

3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

#### LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON


10 model calls with the guard OFF -> 8 with it ON

[Authority comparator] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)

[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Memory provenance firewall

Retrieved long-term memories are testimony, not fact. Each one is checked for origin, signature, staleness and smuggled instructions before it may enter working context - and what gets quarantined is reported, not silently dropped.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                                                                             |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Recall bundle              | Received input - "Paste the memories retrieved for this task, with whatever origin metadata and signatures they carry, plus the current ta" |
| Provenance examiner        | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Content screener           | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                                               |
| Trust adjudicator          | Scored iteration 1/5 - Score: 0.72 - "Weak transitions between key points"                                                                  |
| ... 2 more steps ...       |                                                                                                                                             |
| Trust adjudicator          | Scored iteration 4/5 - Score: 0.98 - "Approved"                                                                                             |
| Admission gate             | Evaluated condition - Condition true -> continue                                                                                            |
| Admitted context + quarant | Published result - Output document ready - see the Output tab                                                                               |


4 critic iteration(s) - final score 0.98 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

11 model calls with the guard OFF -> 7 with it ON  
[Trust adjudicator] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

MCP supply-chain verifier

Before an agent, MCP server or tool joins a graph: identity, publisher, hash, permissions and schema diffed against the last approved version, then exercised in a sandbox so the verdict rests on observed behaviour. Approval is a signed manifest.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                            |                                                                                                                                             |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Component submission       | Received input - "Paste or attach the component proposed for the graph - agent, MCP server, tool or model endpoint: its manifest, publishe" |
| Manifest + identity reader | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Diff against approved      | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                                               |
| Sandbox exerciser          | Writing implementation - 2 files changed - lint clean - 28 tests passing                                                                    |
| ... 2 more steps ...       |                                                                                                                                             |
| Supply-chain adjudicator   | Scored iteration 3/4 - Score: 0.97 - "Approved"                                                                                             |
| Approval gate              | Evaluated condition - Condition true -> continue                                                                                            |
| Signed approved-component  | Published result - Output document ready - see the Output tab                                                                               |


3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

10 model calls with the guard OFF -> 8 with it ON  
[Supply-chain adjudicator] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)

Cost-quality model router

The cheapest capable model goes first and a checker decides if the result actually clears the bar - escalating one rung at a time only on failure. The ledger reports cost per accepted result and savings against an all-frontier baseline, measured, not assumed.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                           |                                                                                                                                             |
|---------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Task batch + model ladder | Received input - "Paste the tasks to run, the model ladder from cheapest local model to frontier (add each rung as a model in Settings and" |
| Task triager              | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                              |
| Cheapest-rung attempt     | Drafting content - 1,428 words - 4 sections - tone matched to brief                                                                         |
| Escalation attempt        | Drafting content - 1,140 words - 5 sections - tone matched to brief                                                                         |
| ... 2 more steps ...      |                                                                                                                                             |
| Acceptance checker        | Scored iteration 3/6 - Score: 0.97 - "Approved"                                                                                             |
| Cost ledger               | Reading and summarizing - 7 documents condensed - key points cross-referenced                                                               |
| Results + routing report  | Published result - Output document ready - see the Output tab                                                                               |

3 critic iteration(s) - final score 0.97 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON


15 model calls with the guard OFF -> 9 with it ON  
[Acceptance checker] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 24 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Meta

Loops that test the loops: hallucination, prompt injection, drift.

Hallucination bench

A researcher paired with a fact-checker that assumes every claim is false until the answer supports it. Break on the step where the citations start being invented.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                          |                                                                                                               |
|--------------------------|---------------------------------------------------------------------------------------------------------------|
| Question                 | Received input - "State the question the researcher should answer, and any sources it is allowed to rely on." |
| Researcher               | Searched 11 sources - 11 sources scanned - 8 relevant - key themes extracted to shared memory                 |
| Adversarial fact-checker | Scored iteration 1/6 - Score: 0.76 - "Claims need stronger sourcing"                                          |
| Adversarial fact-checker | Scored iteration 2/6 - Score: 0.88 - "Weak transitions between key points"                                    |
| Adversarial fact-checker | Scored iteration 3/6 - Score: 0.97 - "Approved"                                                               |
| Claim ledger             | Reading and summarizing - 9 documents condensed - key points cross-referenced                                 |
| Checked answer           | Published result - Output document ready - see the Output tab                                                 |


3 critic iteration(s) - final score 0.97 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 7 with it ON  
[Adversarial fact-checker] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Prompt injection sandbox

An attacker loop hammering an agent with overrides, encoded payloads and indirect injection - with the local Axiom preflight in front, so a blocked attempt costs no model call at all.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                 |                                                                                                                              |
|-----------------|------------------------------------------------------------------------------------------------------------------------------|
| Attack brief    | Received input - "Describe the agent under test, the instructions it is meant to hold to, and what would count as a breach." |
| Attacker        | Drafting content - 1,170 words - 6 sections - tone matched to brief                                                          |
| Axiom preflight | Evaluated condition - Condition true -> continue                                                                             |
| Target agent    | Reading and summarizing - 7 documents condensed - key points cross-referenced                                                |
| Breach detector | Scored iteration 1/6 - Score: 0.71 - "Weak transitions between key points"                                                   |
| Breach detector | Scored iteration 2/6 - Score: 0.82 - "Conclusion does not follow from the analysis"                                          |
| Breach detector | Scored iteration 3/6 - Score: 0.94 - "Approved"                                                                              |
| Red-team report | Published result - Output document ready - see the Output tab                                                                |


3 critic iteration(s) - final score 0.94 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

13 model calls with the guard OFF -> 7 with it ON  
[Breach detector] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Drift and oscillation detector

An optimizer and a scorer, run long enough to misbehave. Watch the trajectory flatten or turn non-monotonic, and watch the governor stop it before the budget goes.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                         |                                                                                              |
|-------------------------|----------------------------------------------------------------------------------------------|
| Seed text + objective   | Received input - "Paste the text to optimize and state precisely what would make it better." |
| Optimizer               | Drafting content - 1,170 words - 6 sections - tone matched to brief                          |
| Scorer                  | Scored iteration 1/12 - Score: 0.76 - "Claims need stronger sourcing"                        |
| Scorer                  | Scored iteration 2/12 - Score: 0.88 - "Weak transitions between key points"                  |
| Scorer                  | Scored iteration 3/12 - Score: 0.97 - "Approved"                                             |
| Final text + trajectory | Published result - Output document ready - see the Output tab                                |


3 critic iteration(s) - final score 0.97 - 2 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

24 model calls with the guard OFF -> 6 with it ON  
[Scorer] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[Trajectory governor] loop node passed context through

Context degradation stress test

Give an agent strict rules, then make it work through ten tasks. The auditor re-checks every answer against the ORIGINAL rules, so you see the iteration where the constraints start slipping rather than discovering it in production.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                   |                                                                                                                                             |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Rules + task list | Received input - "State the rules the agent must hold to for the whole session -- format, forbidden words, length -- then list the tasks i" |
| Task worker       | Drafting content - 1,170 words - 6 sections - tone matched to brief                                                                         |
| Drift auditor     | Scored iteration 1/10 - Score: 0.76 - "Claims need stronger sourcing"                                                                       |
| Drift auditor     | Scored iteration 2/10 - Score: 0.88 - "Weak transitions between key points"                                                                 |
| Drift auditor     | Scored iteration 3/10 - Score: 0.97 - "Weak transitions between key points"                                                                 |
| Drift auditor     | Scored iteration 4/10 - Score: 0.99 - "Approved"                                                                                            |
| Constraint ledger | Reading and summarizing - 5 documents condensed - key points cross-referenced                                                               |
| Drift report      | Published result - Output document ready - see the Output tab                                                                               |


4 critic iteration(s) - final score 0.99 - 3 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

21 model calls with the guard OFF -> 7 with it ON  
[Drift auditor] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 18 check(s), 0 blocked, 7 loop iteration(s) stopped early (~238 estimated tokens saved)

Handoff contagion isolator

When a downstream agent starts hallucinating, the corruption usually arrived from further up. Watches the payload at every handoff, quarantines the node that broke the cascade, and walks backwards to where the context actually went wrong.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                      |                                                                                                                                    |
|----------------------|------------------------------------------------------------------------------------------------------------------------------------|
| Run trace + payloads | Received input - "Paste the run: the sequence of agents, what each one received, what each one produced, and where it went wrong." |
| Handoff interceptor  | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                     |
| Contract checker     | Reading and summarizing - 11 documents condensed - key points cross-referenced                                                     |
| Corruption validator | Scored iteration 1/6 - Score: 0.72 - "Weak transitions between key points"                                                         |
| ... 5 more steps ... |                                                                                                                                    |
| Repro builder        | Writing implementation - 2 files changed - lint clean - 32 tests passing                                                           |
| Cascade report       | Drafting content - 1,233 words - 5 sections - tone matched to brief                                                                |
| Isolated cause       | Published result - Output document ready - see the Output tab                                                                      |


4 critic iteration(s) - final score 0.98 - 6 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

16 model calls with the guard OFF -> 10 with it ON  
[Corruption validator] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 27 check(s), 0 blocked, 3 loop iteration(s) stopped early (~102 estimated tokens saved)

Black-box contract enforcer

Treats every third-party API as a hostile input. Polls the live endpoint, diffs the real response against the schema you rely on, and traps a silent breaking change in a sandbox before it reaches your pipeline.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

**Endpoint + expected schema** Received input - "Give the endpoint, a sample of the response your code expects, and the fields you actually read from it."  
**Live capture** Reading and summarizing - 11 documents condensed - key points cross-referenced  
**Schema differ** Writing implementation - 6 files changed - lint clean - 18 tests passing  
**Sandbox replay** Writing implementation - 2 files changed - lint clean - 28 tests passing  
... 2 more steps ...  
**Breaking or benign?** Scored iteration 3/5 - Score: 0.97 - "Approved"  
**Quarantine before the pipe** Evaluated condition - Condition true -> continue  
**Contract report** Published result - Output document ready - see the Output tab


3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

12 model calls with the guard OFF -> 8 with it ON  
[Breaking or benign?] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 21 check(s), 0 blocked, 2 loop iteration(s) stopped early (~68 estimated tokens saved)

Repo regression bisector

Finds the commit that broke the build when the logs are useless, then does the part bisect cannot: reconstructs what that change was trying to do, maps it to the symptom, and drafts the patch.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

**Good commit, bad commit, f** Received input - "Give a commit where it worked, a commit where it does not, and the test or behaviour that fails."  
**Bisect driver** Writing implementation - 6 files changed - lint clean - 12 tests passing  
**Build and test** Writing implementation - 6 files changed - lint clean - 18 tests passing  
**Narrowed enough?** Scored iteration 1/12 - Score: 0.72 - "Weak transitions between key points"  
... 3 more steps ...  
**Intent reconstructor** Searched 12 sources - 12 sources scanned - 9 relevant - key themes extracted to shared memory  
**Patch proposal** Drafting content - 886 words - 5 sections - tone matched to brief  
**Culprit + patch** Published result - Output document ready - see the Output tab


4 critic iteration(s) - final score 0.98 - 5 agent(s) - run completed

LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON

27 model calls with the guard OFF -> 9 with it ON  
[Narrowed enough?] Loop stopped by Axiom guard - LOOP\_STALLED: draft unchanged for 2 iterations (novelty 0)  
[axiom-guard] 24 check(s), 0 blocked, 9 loop iteration(s) stopped early (~306 estimated tokens saved)

Sarcastic code reviewer

Drop in a messy function. It finds the real bug, hands the finding to an agent with a very different brief, and still gives you the clean fix. The shortest way to see distinct personalities handing off inside one loop.




SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                                                                                                   |                                                                               |
|-------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| Your messy function                                                                                               | Received input - "Paste a function with a bug in it. Messier is funnier."     |
| Bug spotter                                                                                                       | Writing implementation - 6 files changed - lint clean - 12 tests passing      |
| The roast                                                                                                         | Drafting content - 1,428 words - 4 sections - tone matched to brief           |
| The actual fix                                                                                                    | Writing implementation - 2 files changed - lint clean - 28 tests passing      |
| Does it actually work?                                                                                            | Scored iteration 1/3 - Score: 0.72 - "Claims need stronger sourcing"          |
| Does it actually work?                                                                                            | Scored iteration 2/3 - Score: 0.85 - "Good structure, needs tighter language" |
| Does it actually work?                                                                                            | Scored iteration 3/3 - Score: 0.97 - "Approved"                               |
| Fixed code                                                                                                        | Published result - Output document ready - see the Output tab                 |
| 3 critic iteration(s) - final score 0.97 - 4 agent(s) - run completed                                             |                                                                               |
| LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON                                                                    |                                                                               |
| 8 model calls with the guard OFF -> 8 with it ON (nothing to stop; zero overhead)                                 |                                                                               |
| [Does it actually work?] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0) |                                                                               |
| [axiom-guard] 21 check(s), 0 blocked, 0 loop iteration(s) stopped early (~0 estimated tokens saved)               |                                                                               |

Angry customer de-escalator

Paste in a furious all-caps email. The loop separates the heat from the actual request, checks the proposed remedy against policy before a word is drafted, and writes a calm reply that promises only what was authorized.



SIM RUN - REAL ORCHESTRATION, ZERO MODEL CALLS

|                                                                                                                   |                                                                                                                |
|-------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| The email, as received                                                                                            | Received input - "Paste the customer email exactly as it arrived, plus anything relevant about their account." |
| Core issue extractor                                                                                              | Reading and summarizing - 11 documents condensed - key points cross-referenced                                 |
| Policy guard                                                                                                      | Evaluated condition - Condition true -> continue                                                               |
| Draft the reply                                                                                                   | Drafting content - 1,554 words - 3 sections - tone matched to brief                                            |
| Tone and promise check                                                                                            | Scored iteration 1/4 - Score: 0.71 - "Weak transitions between key points"                                     |
| Tone and promise check                                                                                            | Scored iteration 2/4 - Score: 0.82 - "Conclusion does not follow from the analysis"                            |
| Tone and promise check                                                                                            | Scored iteration 3/4 - Score: 0.94 - "Approved"                                                                |
| Reply ready to send                                                                                               | Published result - Output document ready - see the Output tab                                                  |
| 3 critic iteration(s) - final score 0.94 - 3 agent(s) - run completed                                             |                                                                                                                |
| LIVE RUN, STALLED-STUB MODEL - GUARD OFF vs ON                                                                    |                                                                                                                |
| 9 model calls with the guard OFF -> 7 with it ON                                                                  |                                                                                                                |
| [Tone and promise check] Loop stopped by Axiom guard - LOOP_STALLED: draft unchanged for 2 iterations (novelty 0) |                                                                                                                |
| [axiom-guard] 18 check(s), 0 blocked, 1 loop iteration(s) stopped early (~34 estimated tokens saved)              |                                                                                                                |